

Spotify Genre Classification- Statistical Learning Project

Felice Francario

Contents

1	Introduction	1
1.1	Dataset	2
2	Exploratory Data Analysis	4
2.1	Correlation between Features	5
2.2	Feature Distributions	7
3	Models	19
3.1	Multinomial Logistic regression	20
3.2	K-Nearest Neighbours	44
3.3	Random Forest	45
3.4	Gradient Boosting	47

1 Introduction

The objective of this project is to analyze the diverse landscape of music genres and develop a classification model capable of predicting the genre of a song based on its audio features. Music, as a universal form of expression, encompasses a vast array of genres, each characterized by distinct musical elements, themes, and emotions. Understanding and classifying these genres is not only valuable for music enthusiasts and artists but also holds significant implications for music recommendation systems, content curation, and genre-specific analysis.

The proliferation of digital music platforms and streaming services has led to an exponential growth in the availability and accessibility of music across genres. However, manually categorizing and organizing this vast musical repertoire is a daunting task. Automated classification models offer a promising solution by leveraging machine learning algorithms to discern patterns and relationships among audio features and genre labels.

In this project, we embark on a comprehensive exploration of music genres, drawing insights from a dataset of songs spanning diverse genres and cultural influences. By analyzing key audio features such as tempo, energy, danceability, and instrumentality, we aim to uncover the underlying characteristics that distinguish one genre from another.

Our primary goal is to develop a robust classification model capable of accurately predicting the genre of a song based solely on its audio features.

1.1 Dataset

The dataset was downloaded from Kaggle in a CSV format. This dataset provides comprehensive information about Spotify tracks encompassing a diverse collection of 125 genres.

```
data<-read.csv("train.csv")
dim(data)
```

```
## [1] 114000    21
```

The dataset has already been cleaned and contains no missing values

```
sum(is.na(data))
```

```
## [1] 0
```

The dataset comprises multiple columns, each representing distinctive audio features associated with individual tracks.

```
names<-colnames(data)
names
```

```
## [1] "Unnamed..0"      "track_id"        "artists"         "album_name"
## [5] "track_name"      "popularity"      "duration_ms"     "explicit"
## [9] "danceability"    "energy"          "key"             "loudness"
## [13] "mode"            "speechiness"     "acousticness"    "instrumentalness"
## [17] "liveness"        "valence"         "tempo"           "time_signature"
## [21] "track_genre"
```

The first 5 descriptive columns were removed as they are irrelevant for the objective of the project

```
#removing unnamed..0, track_id, artists, album_name, track_name columns
df<-data[,-c(1,2,3,4,5)]
colnames(df)
```

```
## [1] "popularity"      "duration_ms"     "explicit"        "danceability"
## [5] "energy"          "key"             "loudness"        "mode"
## [9] "speechiness"     "acousticness"    "instrumentalness" "liveness"
## [13] "valence"         "tempo"           "time_signature"  "track_genre"
```

1.1.1 Description of the features

1. **popularity:** A score indicating how popular a track is on Spotify (ranging from 0 to 100).
2. **duration_ms:** The duration of a track in milliseconds.
3. **explicit:** Indicates whether a track contains explicit content (True or False). Explore Audio Features: This dataset includes various audio features associated with each track. Here are some notable ones:
4. **Danceability:** Danceability measures how suitable a track is for dancing, ranging from 0 to 1. Tracks with high danceability scores are more energetic and rhythmic, making them ideal for dancing.

5. **Energy:** Energy represents intensity and activity within a song on a scale from 0 to 1. Tracks with high energy tend to be more fast-paced and intense.
6. **Loudness:** Loudness indicates how loud or quiet an entire song is in decibels (dB). Positive values represent louder songs while negative values suggest quieter ones.
7. **Key:** Key refers to different musical keys assigned integers ranging from 0-11, with each number representing a different key. Knowing the key can provide insights into the mood and tone of a song.
8. **mode:** The tonal mode of the track, represented by an integer value (0 for minor, 1 for major).
9. **speechiness:** A score ranging from 0 to 1 that represents the presence of spoken words in a track.
10. **acousticness:** A score ranging from 0 to 1 that represents the extent to which a track possesses an acoustic quality.
11. **instrumentalness:** A score ranging from 0 to 1 that represents the likelihood of a track being instrumental.
12. **liveness:** A score ranging from 0 to 1 that represents the presence of an audience during the recording or performance of a track.
13. **Valence:** Valence measures the musical positiveness conveyed by a track, ranging from 0 to 1. High valence values indicate more positive or happy tracks, while lower values suggest more negative or sad ones.
14. **Tempo:** Tempo is the speed or pace of a song in beats per minute (BPM). It gives an idea about how fast or slow a track is.
15. **time_signature:** The number of beats within each bar of the track.

1.1.2 Response variable

16. **track_genre:** The genre of the track. This will be the **response variable**.

```
unique(df$track_genre)
```

```
##      [1] "acoustic"          "afrobeat"          "alt-rock"
##      [4] "alternative"       "ambient"           "anime"
##      [7] "black-metal"       "bluegrass"         "blues"
##     [10] "brazil"            "breakbeat"         "british"
##     [13] "cantopop"          "chicago-house"    "children"
##     [16] "chill"             "classical"         "club"
##     [19] "comedy"            "country"           "dance"
##     [22] "dancehall"         "death-metal"       "deep-house"
##     [25] "detroit-techno"    "disco"             "disney"
##     [28] "drum-and-bass"     "dub"               "dubstep"
##     [31] "edm"               "electro"           "electronic"
##     [34] "emo"               "folk"              "forro"
##     [37] "french"            "funk"              "garage"
##     [40] "german"            "gospel"            "goth"
##     [43] "grindcore"         "groove"            "grunge"
##     [46] "guitar"            "happy"             "hard-rock"
##     [49] "hardcore"          "hardstyle"         "heavy-metal"
##     [52] "hip-hop"           "honky-tonk"        "house"
##     [55] "idm"               "indian"            "indie-pop"
```

```
## [58] "indie"          "industrial"      "iranian"
## [61] "j-dance"        "j-idol"          "j-pop"
## [64] "j-rock"         "jazz"            "k-pop"
## [67] "kids"           "latin"           "latino"
## [70] "malay"          "mandopop"        "metal"
## [73] "metalcore"      "minimal-techno"  "mpb"
## [76] "new-age"        "opera"           "pagode"
## [79] "party"          "piano"           "pop-film"
## [82] "pop"            "power-pop"       "progressive-house"
## [85] "psych-rock"     "punk-rock"       "punk"
## [88] "r-n-b"          "reggae"          "reggaeton"
## [91] "rock-n-roll"    "rock"            "rockabilly"
## [94] "romance"        "sad"             "salsa"
## [97] "samba"          "sertanejo"       "show-tunes"
## [100] "singer-songwriter" "ska"            "sleep"
## [103] "songwriter"     "soul"            "spanish"
## [106] "study"          "swedish"         "synth-pop"
## [109] "tango"          "techno"          "trance"
## [112] "trip-hop"       "turkish"         "world-music"
```

Since there are too many different genres, only the 9 most important genres were selected for this analysis.

```
genres<-c('classical','country','electronic','hip-hop','jazz','rock','pop','blues','reggae')

#Filter the dataframe with selected genres
df_filtered<-df[df$track_genre %in% genres,]

y<-as.factor(df_filtered$track_genre)

df_filtered$track_genre <- as.factor(df_filtered$track_genre)
unique(y)
```

```
## [1] blues      classical  country   electronic hip-hop    jazz      pop
## [8] reggae     rock
## Levels: blues classical country electronic hip-hop jazz pop reggae rock
```

The dataset contains exactly 1000 tracks for each genre.

```
table(y)
```

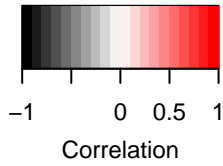
```
## y
##    blues  classical    country electronic    hip-hop    jazz    pop
##    1000    1000      1000      1000      1000      1000  1000
##   reggae     rock
##    1000     1000
```

2 Exploratory Data Analysis

In this section, we conduct exploratory data analysis (EDA) to gain insights into the distributional characteristics of various features in our dataset. EDA plays a crucial role in understanding the underlying patterns and relationships among different variables before proceeding with further analysis or modeling. We will view the relationship between the various features and the distribution of the various features based on the track genre, our response variable.

2.1 Correlation between Features

```
#considering only numeric features
numeric_df <- df_filtered[, sapply(df_filtered, is.numeric)]
cor_matrix<-cor(numeric_df)
variable_names <- colnames(cor_matrix)
heatmap.2(cor_matrix,
  trace = "none", # Turn off row/column dendrogram
  col = colorRampPalette(c("black", "white", "red"))(20),
  main = "Correlation Matrix",
  dendrogram="none",
  #xlab = "Variables", ylab = "Variables",
  key = TRUE, # Add color key for the gradient
  key.title = NA, # Remove the default key title
  key.xlab = "Correlation", # Add x-axis label for the key
  key.ylab = NULL, # Remove the default y-axis label for the key
  Rowv = variable_names, # Set row order
  Colv = variable_names, # Set column order
  density.info = "none", # Turn off density plot
  symkey = FALSE, # Do not plot a symmetric key
  add.expr = {
    # Add text annotations for correlation values
    for (i in 1:nrow(cor_matrix)) {
      for (j in 1:ncol(cor_matrix)) {
        text(i,15-j, format(cor_matrix[i, j], digits = 2),
          col = "black", cex = 0.8)
      }
    }
  },
  #width = 6, # Adjust width of the plot
  #height = 6, # Adjust height of the plot
  cexRow = 0.8, # Adjust text size for rows
  cexCol = 0.8 # Adjust text size for columns
)
```



Correlation Matrix

1	0.11	0.15	0.18	0.00	0.17	-0.15	0.07	-0.19	0.07	0.00	0.48	0.06	0.25	0.07	popularity
0.11	1	-0.08	0.05	-0.01	0.00	0.88	0.03	0.02	0.03	0.00	0.66	-0.17	-0.02	0.00	duration_ms
0.15	-0.08	1	0.44	0.08	0.48	-0.14	0.28	-0.47	-0.31	-0.02	0.44	-0.03	0.50	0.23	danceability
0.18	0.05	0.44	1	0.1	0.82	-0.12	0.22	-0.79	-0.42	0.14	0.37	0.21	0.25		energy
0.00	0.01	0.08	0.1	1	0.08	-0.11	0.06	-0.06	-0.04	0.03	0.50	0.09	0.18	0.02	key
0.17	0.00	0.48	0.82	0.08	1	-0.09	0.18	-0.69	-0.62	0.09	0.3	0.16	0.18		loudness
-0.15	-0.03	-0.14	-0.12	-0.11	-0.09	1	-0.11	0.16	0.01	-0.01	0.14	0.00	0.80	0.02	mode
0.07	0.02	0.28	0.22	0.06	0.18	-0.11	1	-0.17	-0.12	0.11	0.13	0.13	0.08		speechiness
-0.19	-0.03	-0.47	-0.79	-0.06	-0.69	-0.17	1	0.37	-0.05	-0.21	-0.19	-0.23			acousticness
-0.07	0.03	-0.31	-0.42	-0.04	-0.62	0.01	-0.12	0.37	1	-0.05	-0.28	-0.1	-0.09		instrumentalness
0.00	0.00	0.60	0.27	0.14	0.03	0.09	-0.01	0.11	-0.05	0.05	1	0.07	0.80	0.04	liveness
-0.06	-0.17	0.44	0.37	0.09	0.3	0.01	0.13	-0.21	-0.28	0.07	1	0.13	0.15		valence
0.02	-0.02	-0.03	0.21	0.01	0.16	-0.00	0.13	-0.19	-0.1	0.04	0.13	1	0.00		tempo
0.07	0.00	0.23	0.25	0.02	0.18	-0.02	0.08	-0.23	-0.09	0.07	0.10	0.15	0.00	1	time_signature
popularity	duration_ms	danceability	energy	key	loudness	mode	speechiness	acousticness	instrumentalness	liveness	valence	tempo	time_signature		

From the correlation matrix, we can notice that danceability, energy and loudness are significantly correlated negatively with acousticness, instrumentalness while they are positively correlated with each other. In particular there is a very high correlation of 0.82 between energy and loudness.

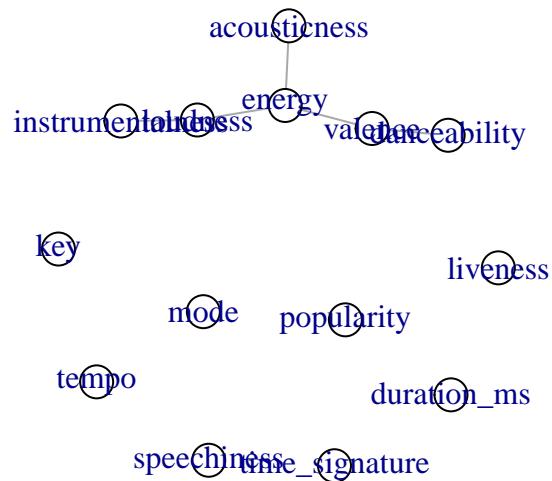
From the following you can see how above mentioned features that are partially correlated with each other while the rest of the features are basically conditionally independent with respect to each other.

```
S <- var(numeric_df)
R <- -cov2cor(solve(S))

thr <- 0.3

G <- abs(R) > thr
diag(G) <- 0

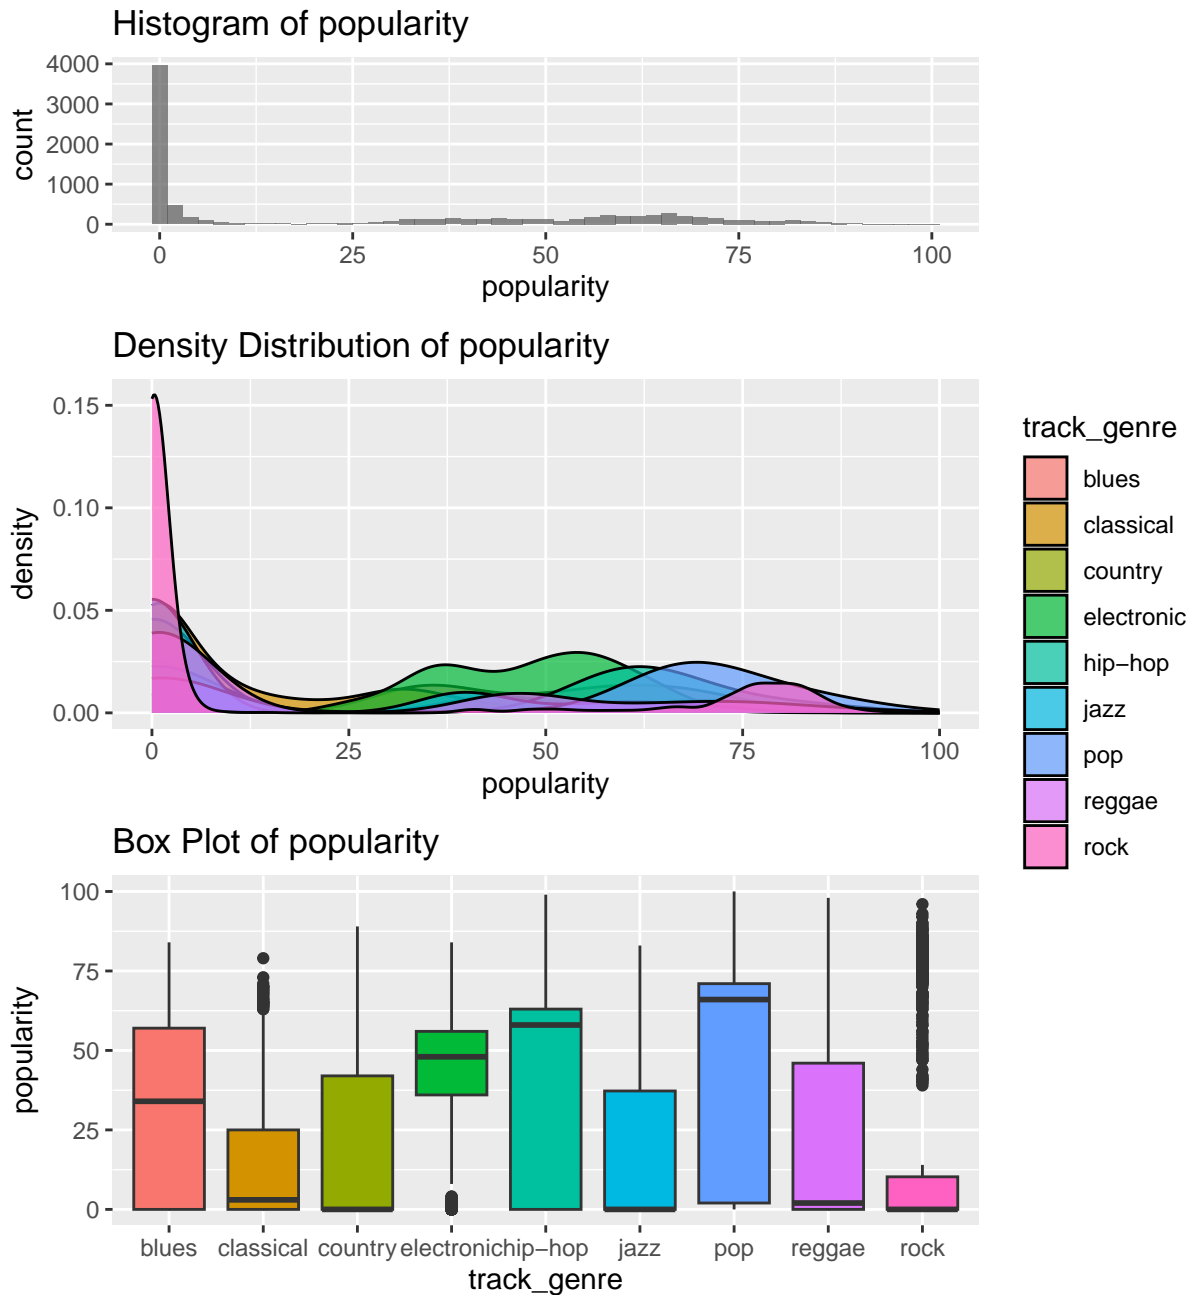
#Gi <- as(G, "igraph")
#plot(Gi)
Gi <- graph_from_adjacency_matrix(as.matrix(G), mode = "undirected", weighted = NULL, diag = FALSE)
plot(Gi, vertex.color = "white")
```



2.2 Feature Distributions

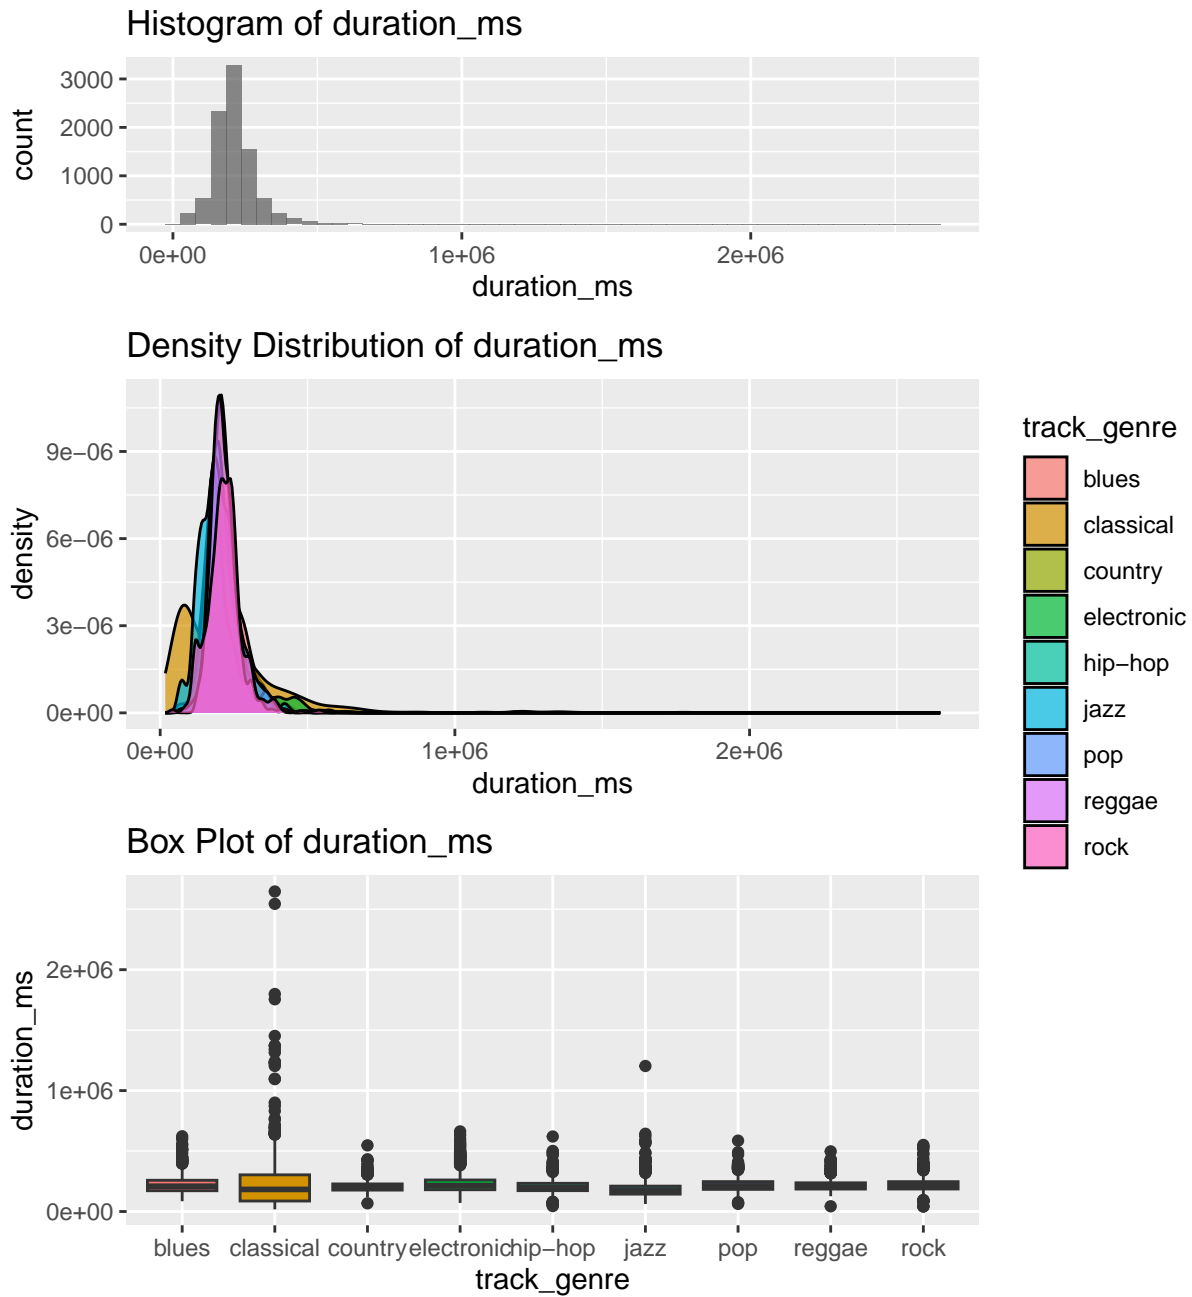
In this subsection we examine the distribution of individual features and their relationship with the target variable. For each feature, we generate histograms, density plots, and box plots to visualize the distribution across different classes.

1. Popularity



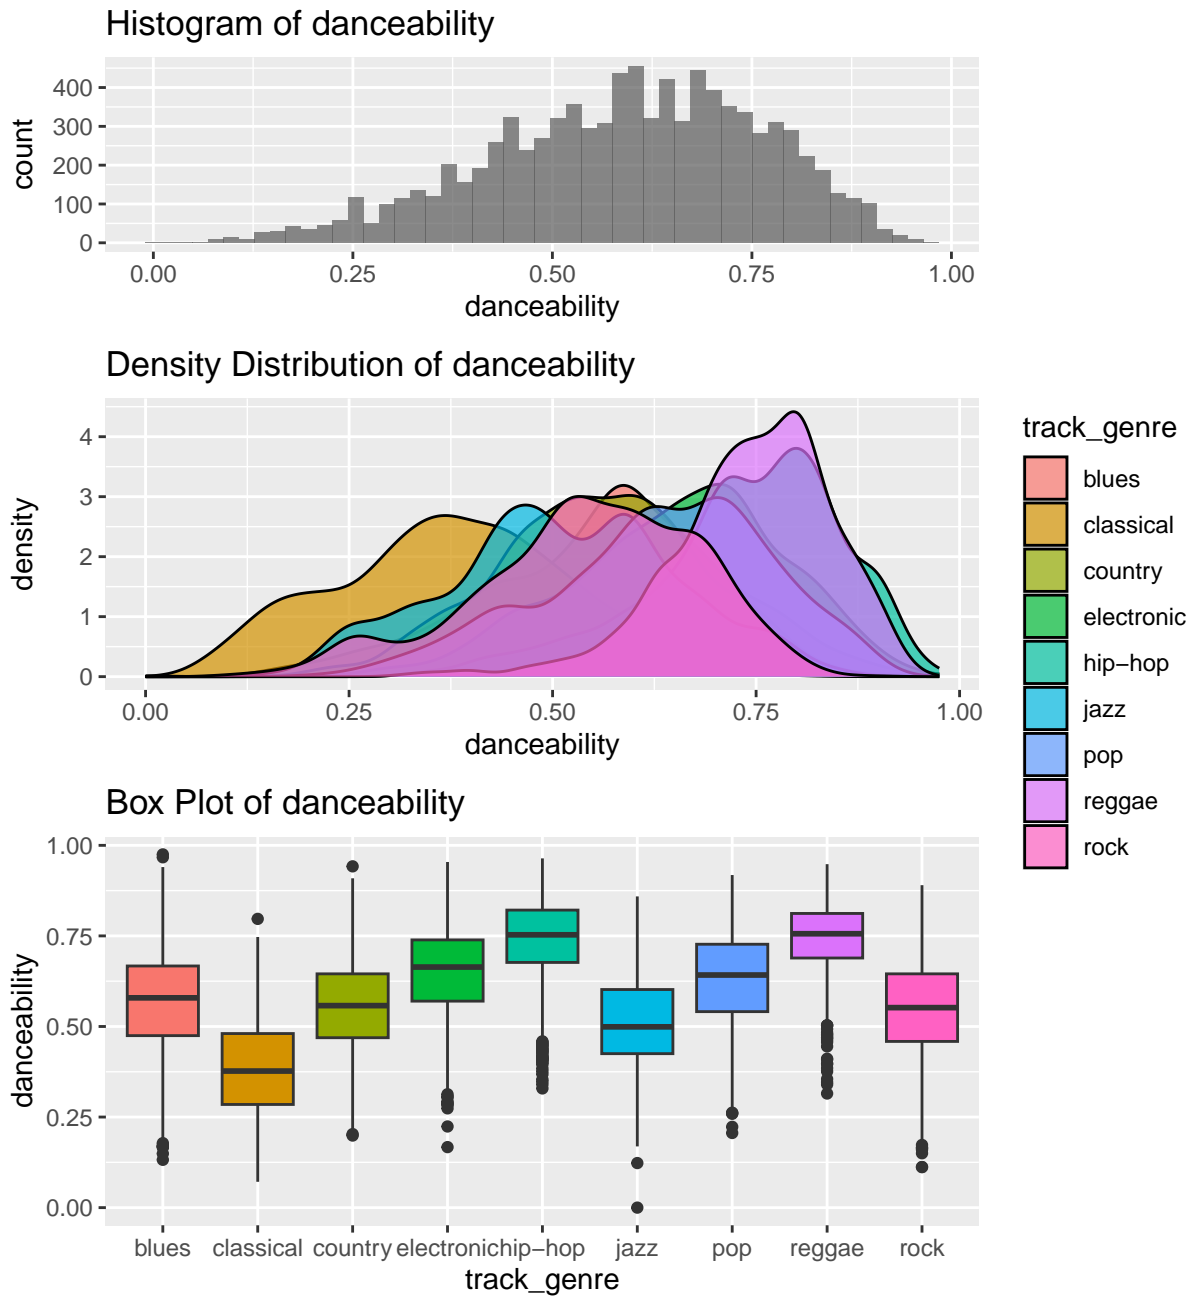
Classical, country, jazz, reggae and rock seem to generally be the least popular genres, with the majority of the rock tracks with a particularly low popularity score, but with a lot of outliers with a very high score. The most popular genres are pop, hip-hop and electronic. Hip-hop and pop have the highest spread in popularity values, with a very high median value but left skewed distributions.

2. Duration-ms



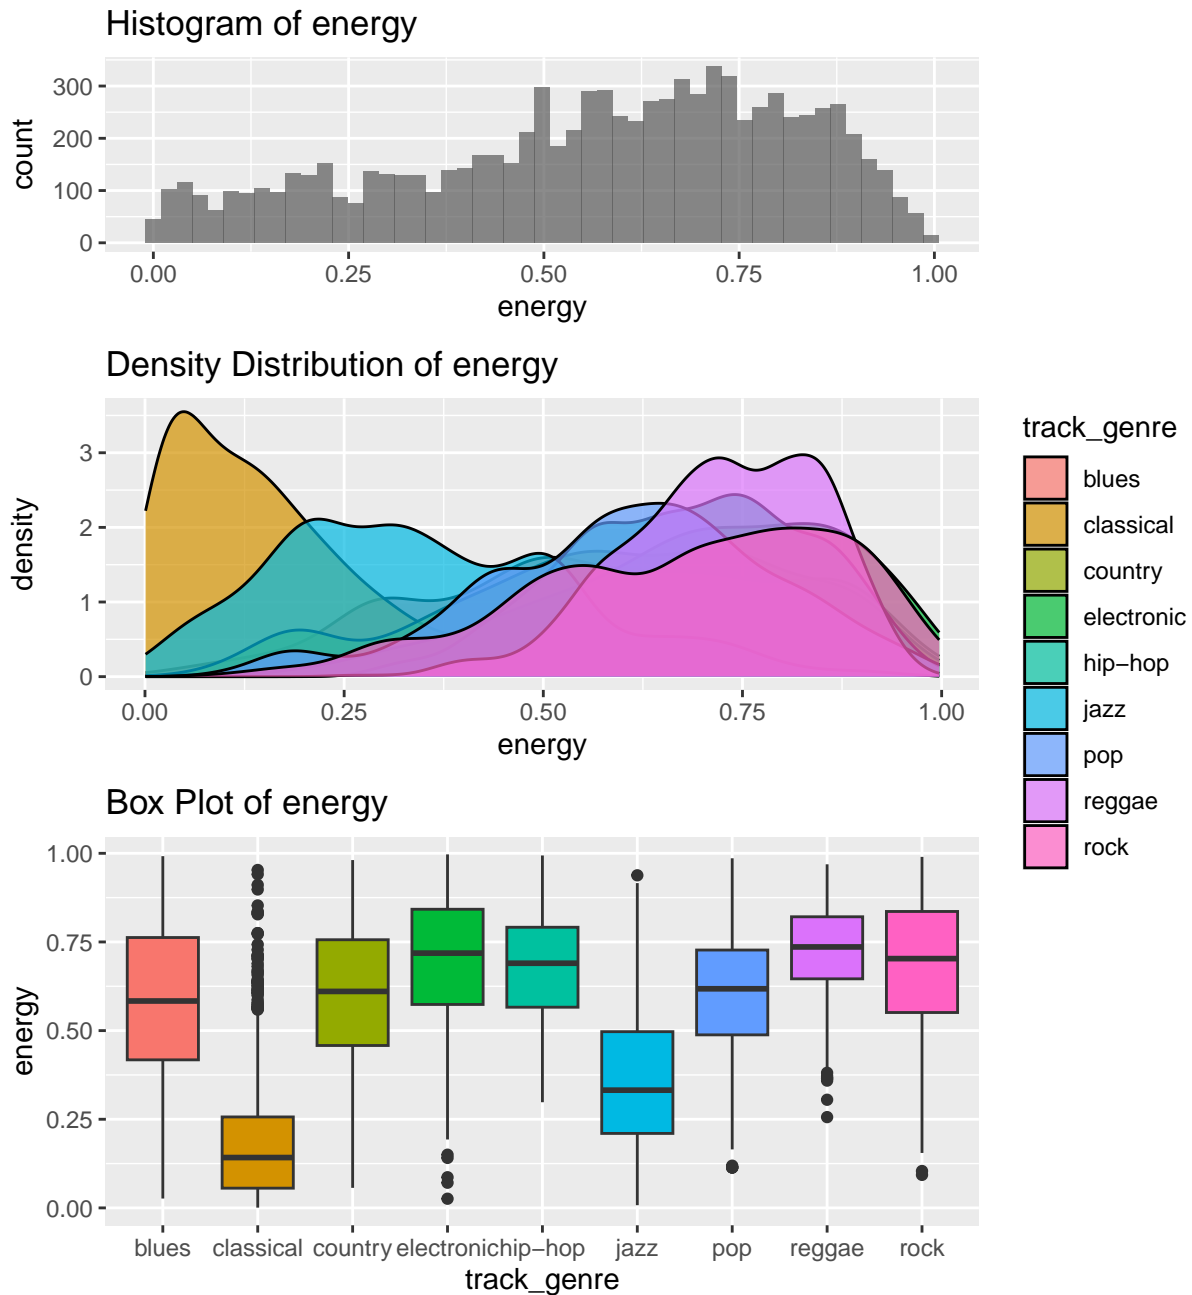
All the genres have a relatively similar distribution except the classical genre which has the highest variance and lowest median value. The classical genre has the highest range of track duration with a relatively high concentration of short tracks but with a huge amount of outliers with extremely high track duration.

3. Danceability



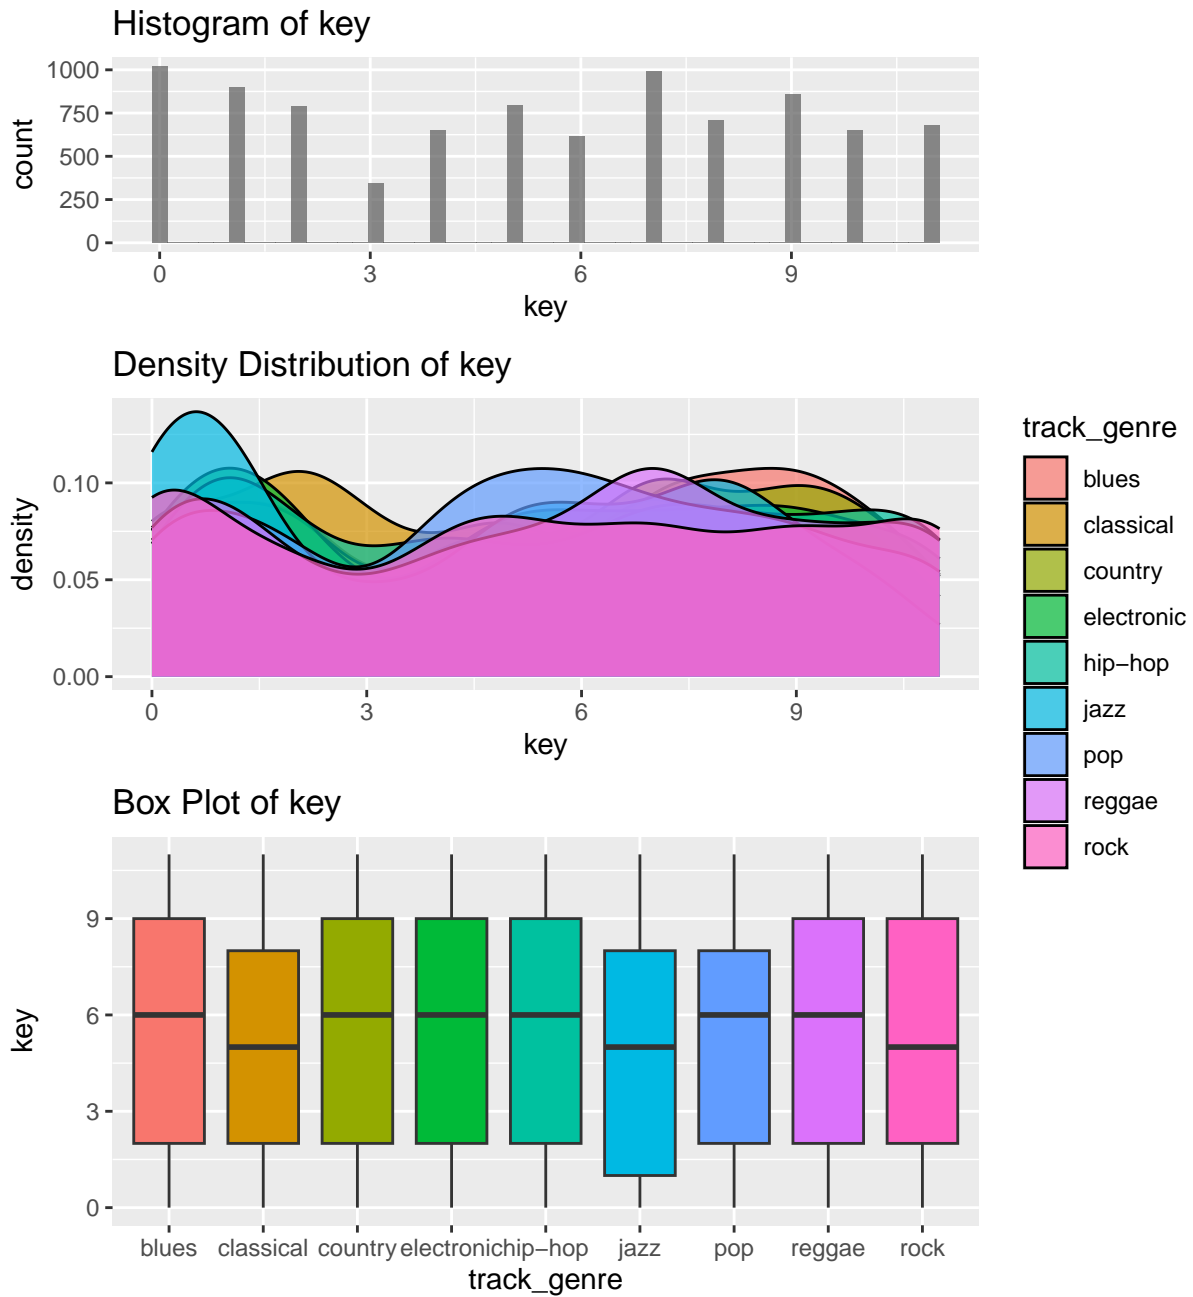
Hip-Hop and Reggae have the heighest danceability scores, with Hip-Hop in general having a slightly higher concentration of tracks near the very maximum value of danceability. Classical and Jazz obviously have the lowest danceability score.

4. Energy



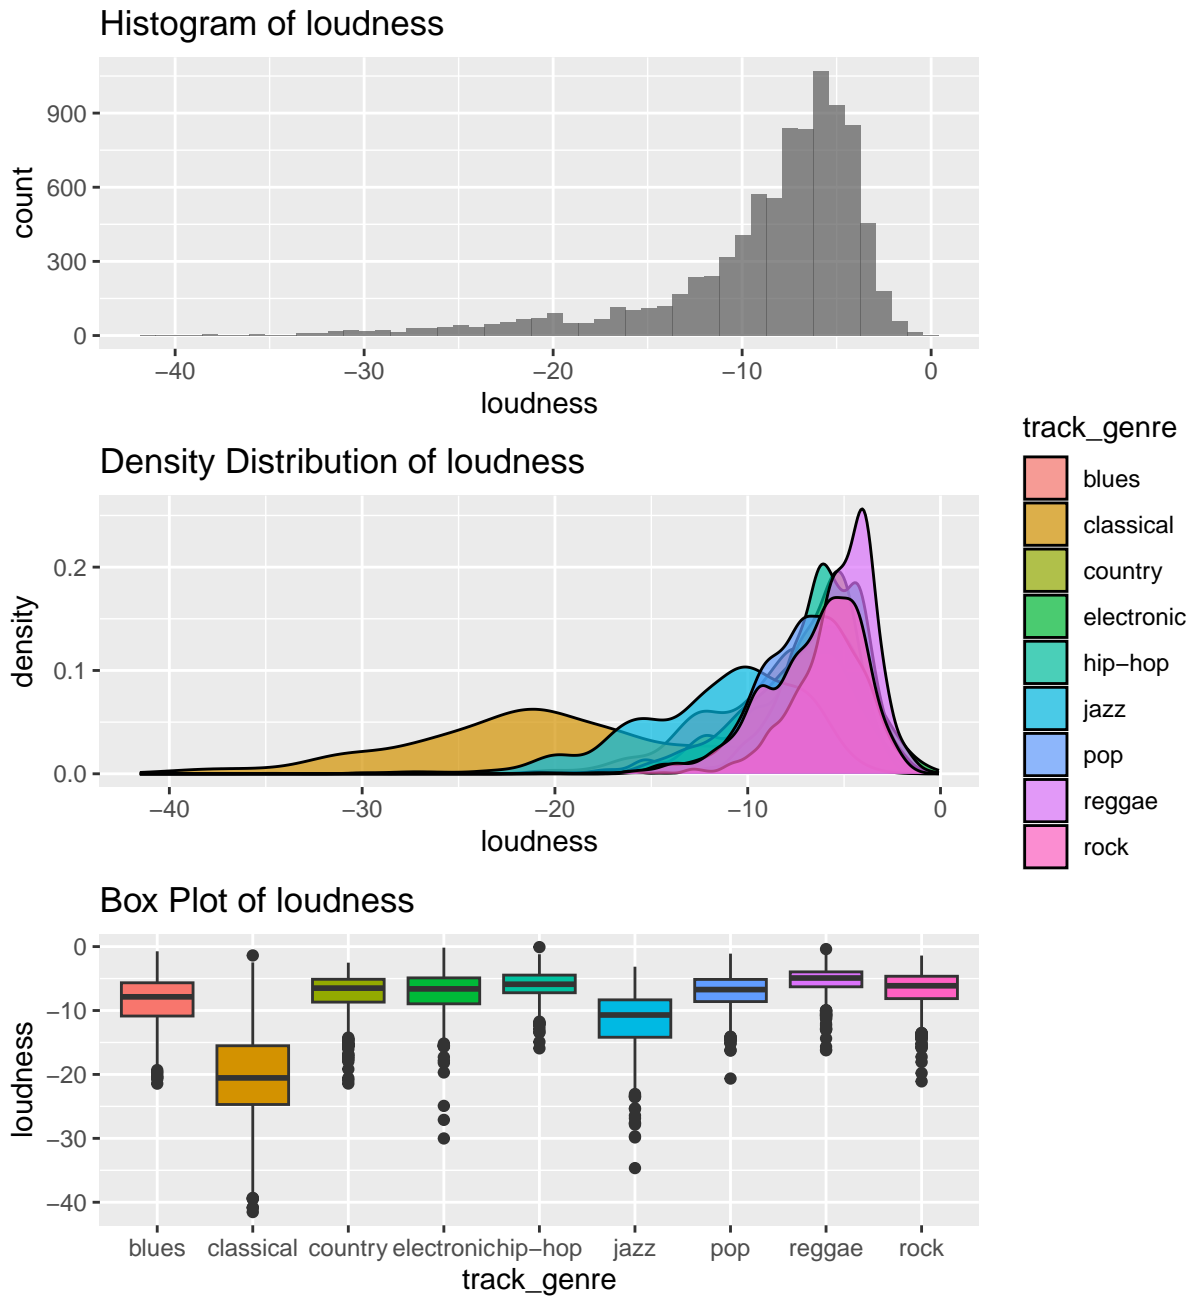
The distributions of the energy score are generally similar to danceability but with even lower scores for classical and jazz tracks. The main differences are that electronic tracks have a higher energy score than hip-hop and rock has a generally much higher energy score compared to danceability.

6. Key



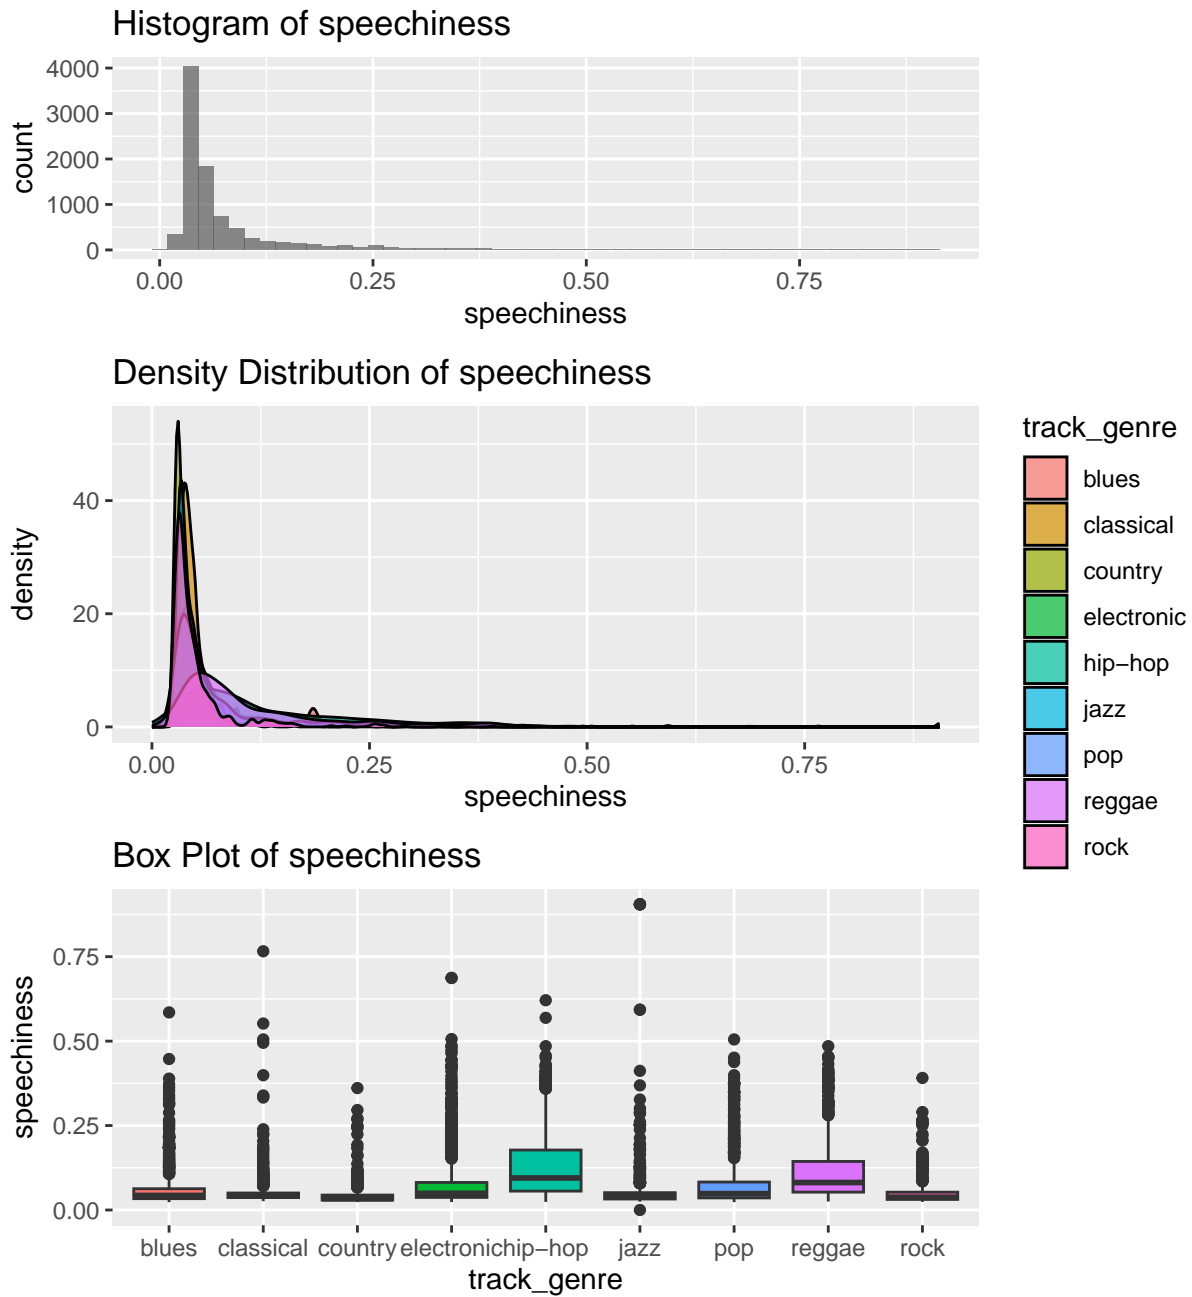
Visually the feature “Key” doesn’t have a big effect on the differences between the genres.

7. Loudness



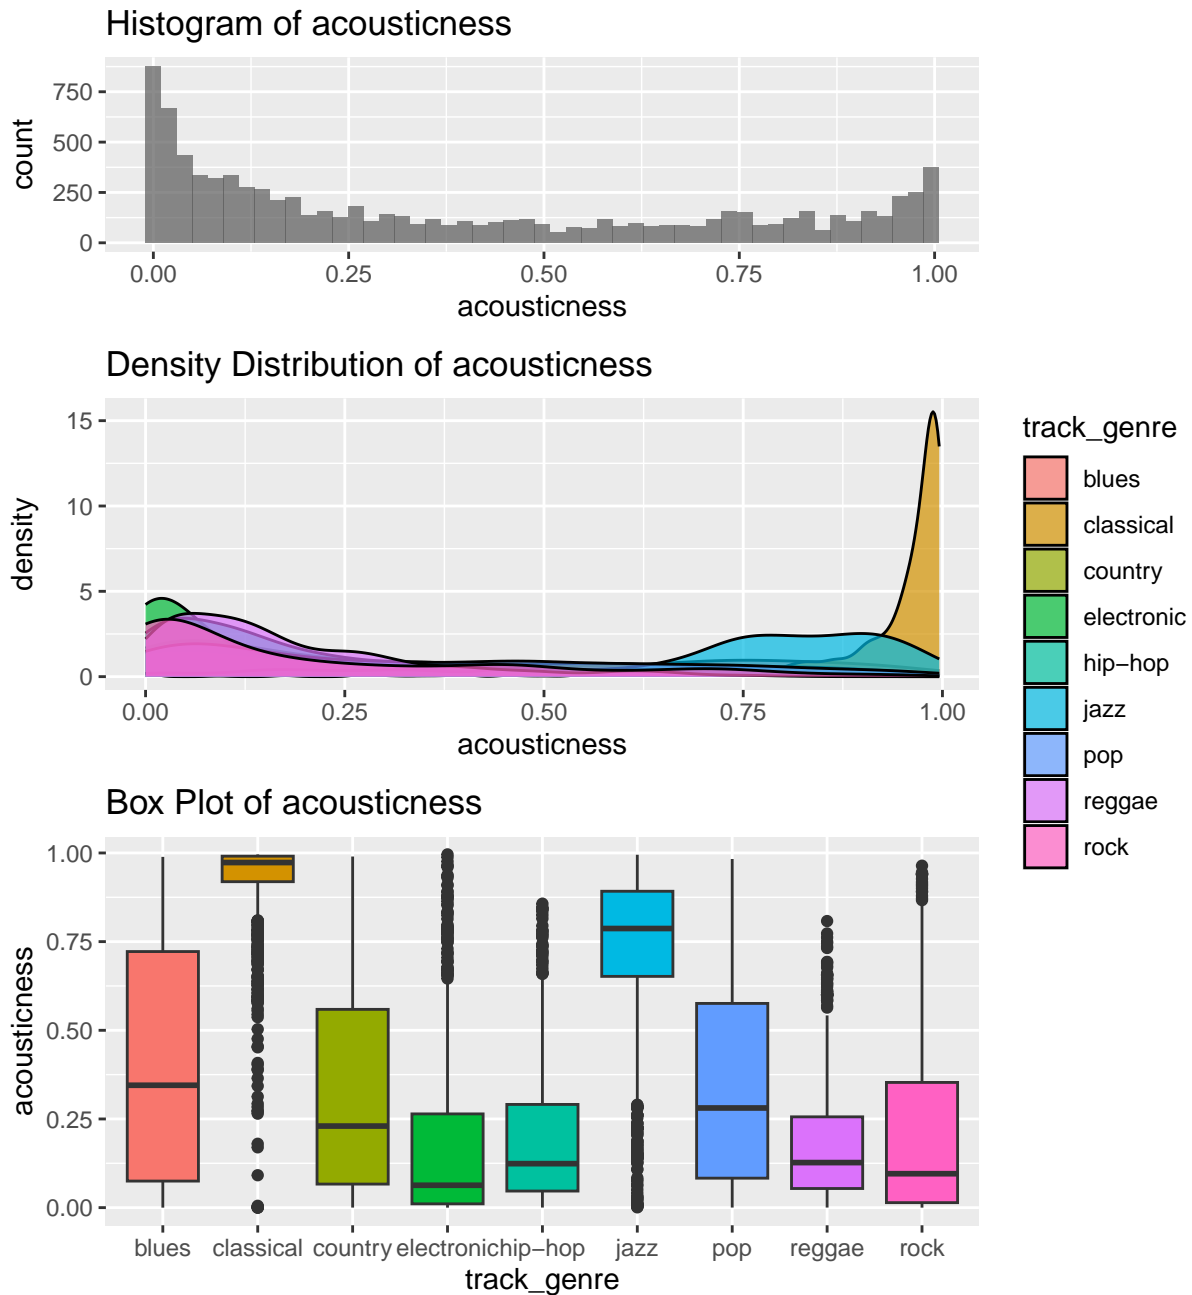
Reggae and Hip-hop are generally the loudest, while jazz and classical are generally quieter. This feature has a 0.82 correlation with energy, so the similarity in distributions is expected.

8. Speechiness



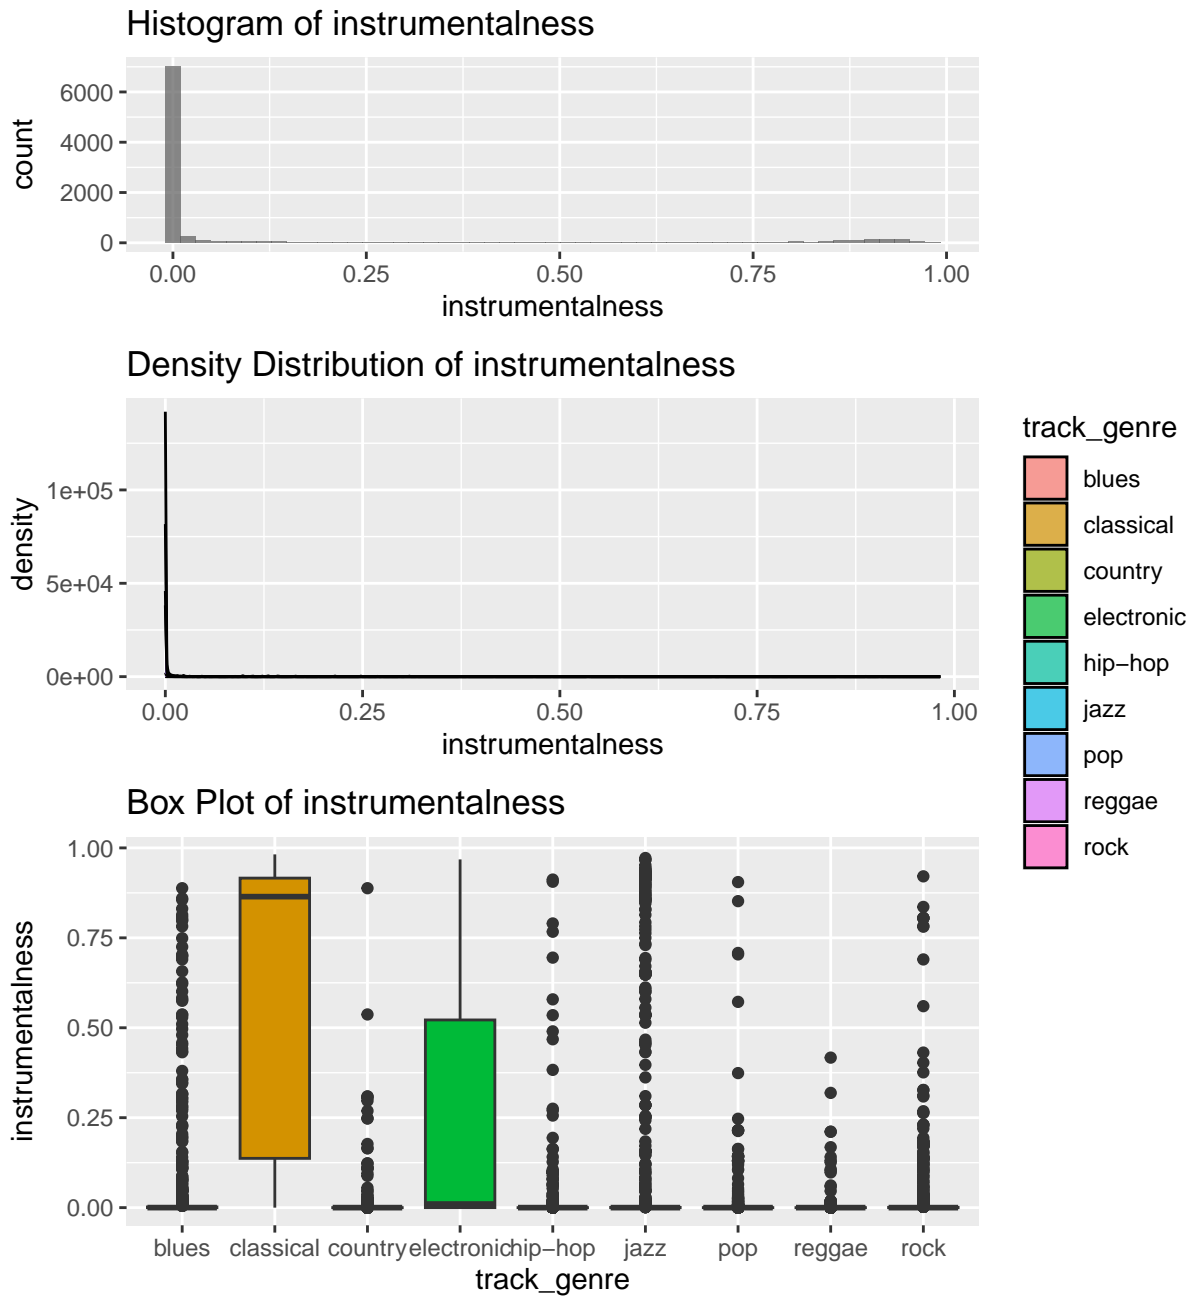
All the genres have a majority of their tracks concentrated at a very low score, with country in particular having the lowest variance while Hip-hop has the highest median value and variance.

9. Acousticness



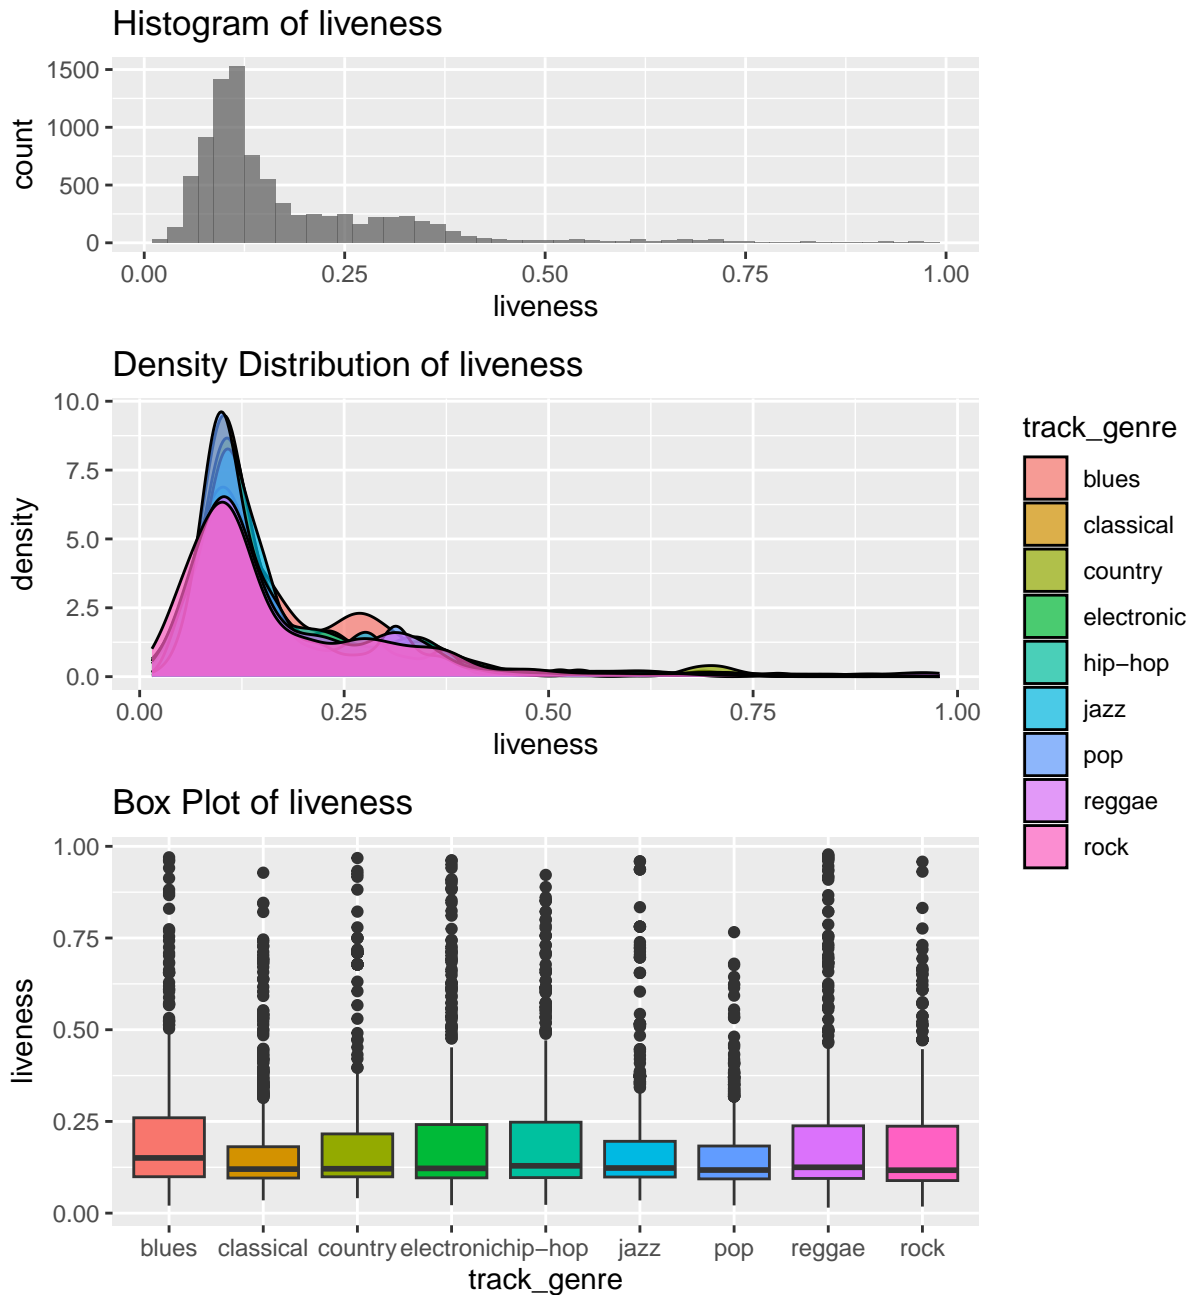
This feature has a high variance for all the genres with the exception of the classical genre which is highly concentrated close to the maximum value of 1. Electronic, hip-hop and reggae have in general the lowest acousticness score.

10. Instrumentalness



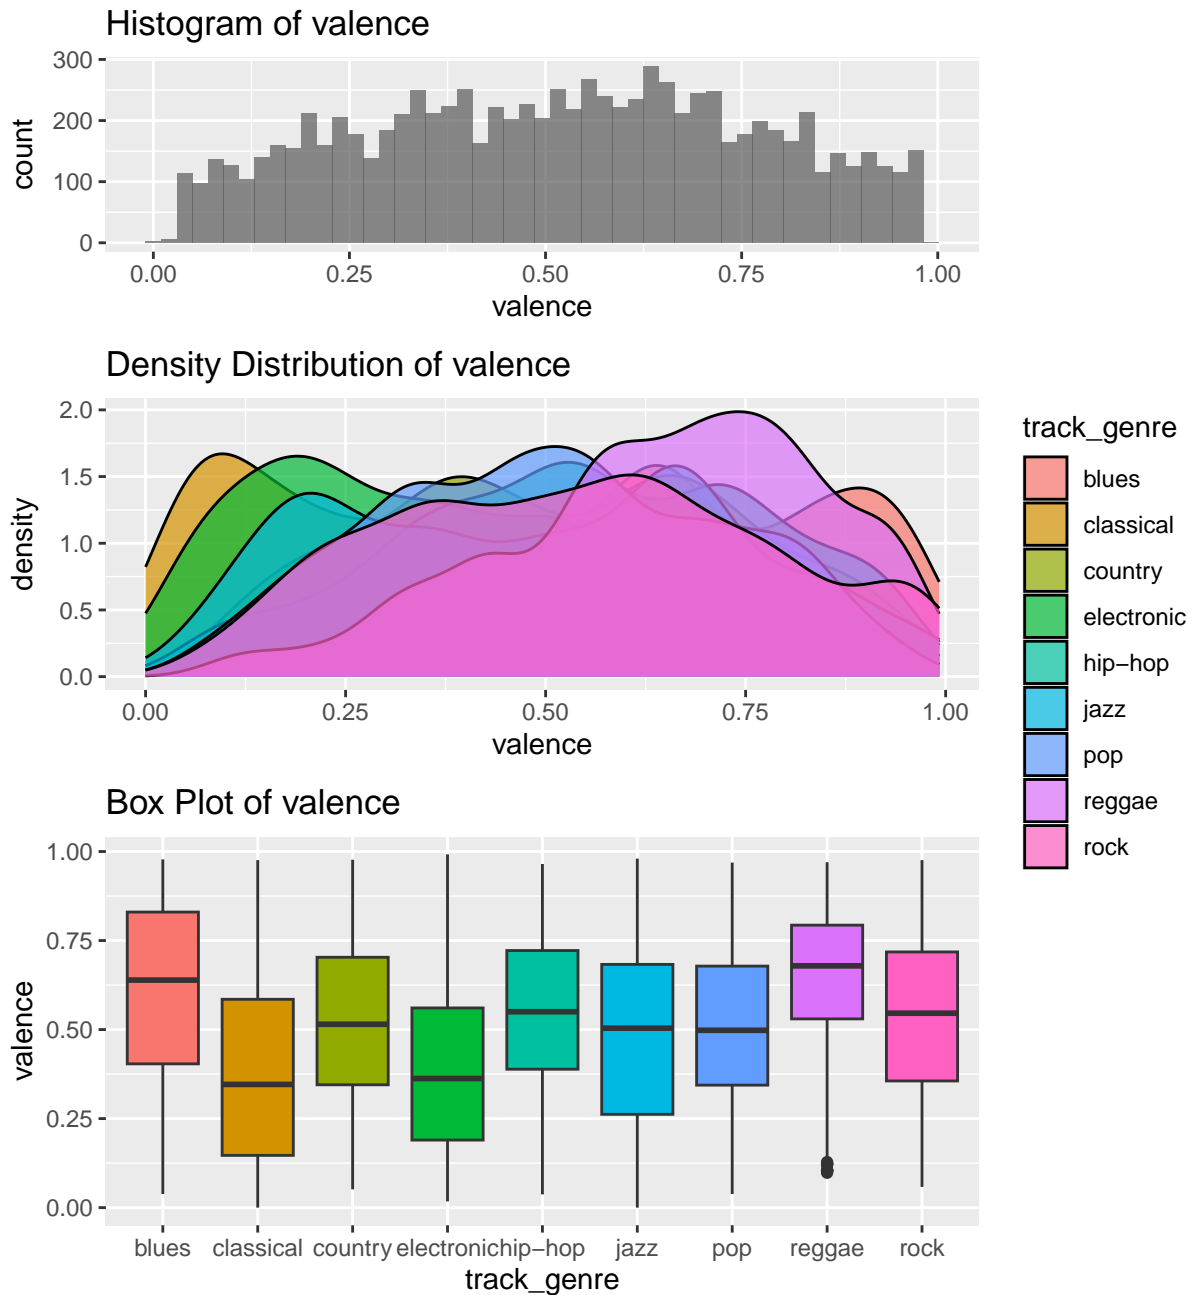
Aside from classical and electronic, all the genres have an instrumentality score close to 0. Classical is heavily left skewed distribution with a high variance while electronic has a very right skewed distribution with its median value close to 0 but a very high variance.

11. Liveness



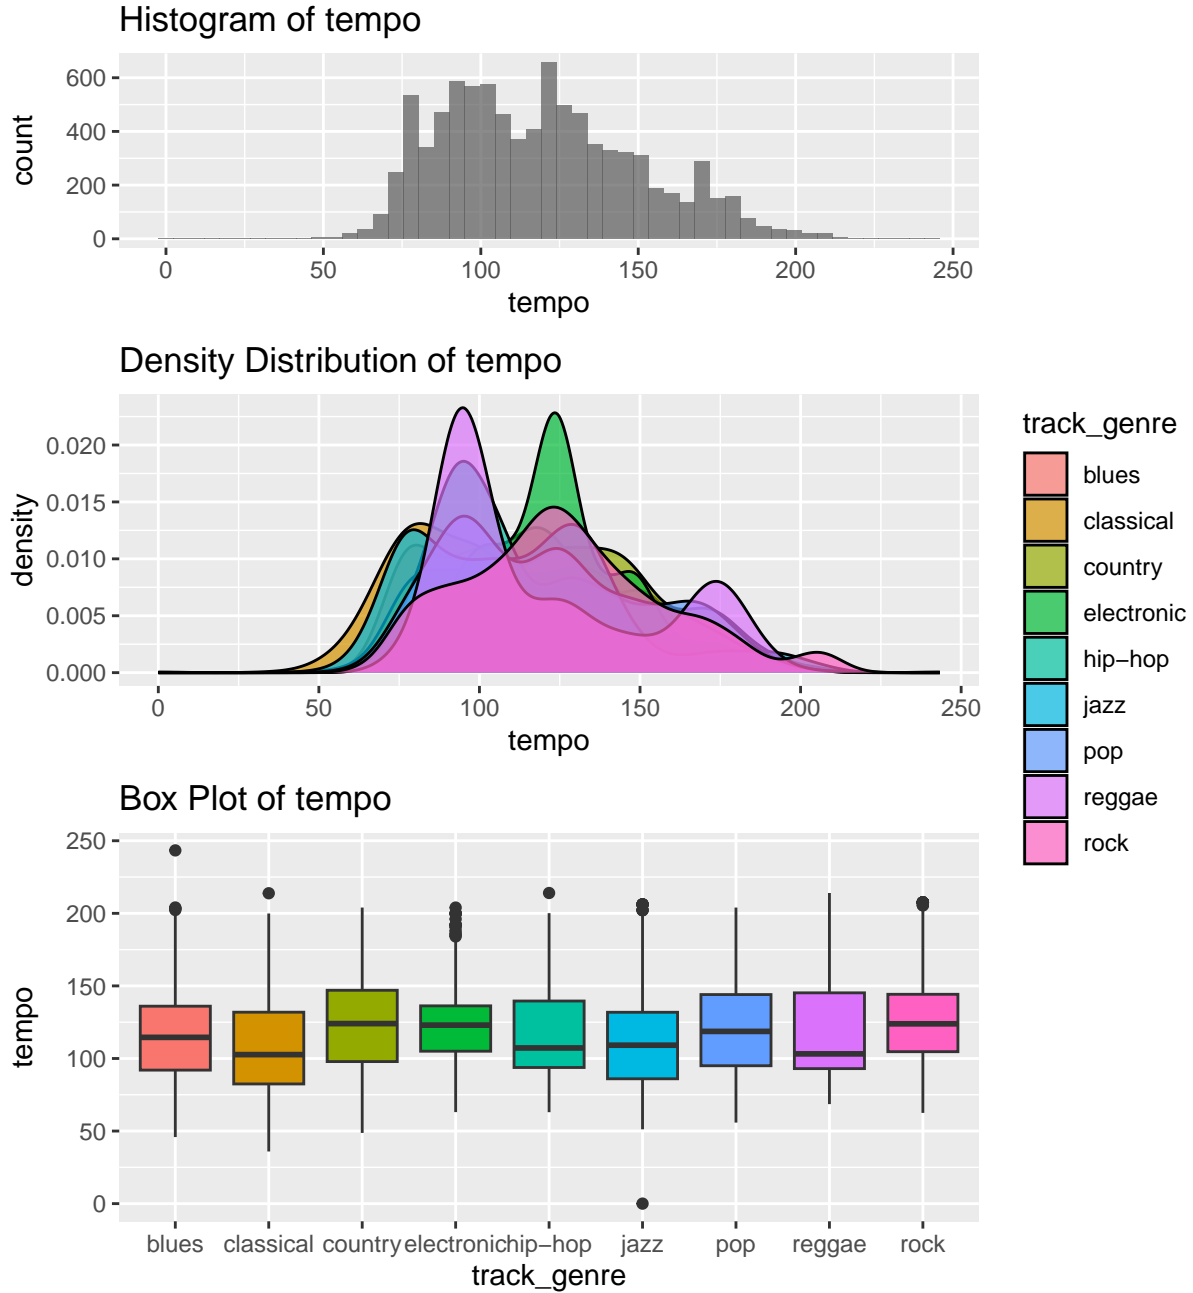
Visually there doesn't seem to be much difference in the liveness distribution of all the genres. They all have a right skewed distribution with median around 0.12 with the exception of the blues genres which has a slightly higher median value than all the other genres. There a lot of outliers for each genre.

12. Valence



Reggae and blues have the highest valence scores, while classical and electronic have the lowest. This indicated that reggae and blues generally have a more positive vibe while classical and electronic are generally more negative.

13. Tempo



The genre distribution are multimodal and have peaks around distinct values. In particular Reggae and electronic have the most identifiable distributions.

3 Models

In this section, we will apply various tuned models trained on a training set for a multi class classification task evaluated on a separate test set. To create the training and test set, the dataset was divided with 80/20 ratio. Here we create a model matrix X and scale it to a normal distribution $N(0,1)$ to reduce the effect of the scale of a feature on predictions.

```

set.seed(1)
X<-model.matrix(track_genre~.,data=df_filtered)[,-1]

numeric_vars <- sapply(df_filtered[,-16], is.numeric)
categorical_vars <- !numeric_vars

# Scale numerical variables in the model matrix
X_scaled<-X
X_scaled[, numeric_vars] <- scale(X[, numeric_vars])

train <- sample(1:nrow(X), 4*nrow(X)/5)
test <- (-train)
#X.train<-X[train,]
#X.test<-X[test,]
y.train<-y[train]
y.test <- y[test]

X_scaled.train<-X_scaled[train,]
X_scaled.test<-X_scaled[test,]

```

3.1 Multinomial Logistic regression

In this subsection we will apply different variants of a Multinomial regression with feature selection or regularization. Multinomial logistic regression is a powerful technique for modeling relationships between multiple categorical outcome variables and predictor variables. The multinomial logistic regression model extends the binary logistic regression model to handle multiple classes. The conditional probability of observing class k given predictor variables \mathbf{x} is given by:

$$P(Y = k|\mathbf{x}) = \frac{e^{\beta_{0k} + \beta_k^T \mathbf{x}}}{1 + \sum_{l=1}^{K-1} e^{\beta_{0l} + \beta_l^T \mathbf{x}}}$$

where K is the number of classes, β_{0k} are the intercepts for each class, β_k are the coefficient vectors for each class, and \mathbf{x} is the vector of predictor variables.

Next, we discuss the implementation details of multinomial logistic regression in R, leveraging popular libraries such as `glmnet` or `nnet`. In particular, we explore the `multinom()` function provided by the `nnet` package in R. The `multinom()` function fits multinomial logistic regression models using the maximum likelihood estimation method and can handle both nominal and ordinal response variables.

3.1.1 Base Model

```

multinom_model <- multinom(track_genre ~ ., data = df_filtered[train,])
#summary(multinom_model)

predictions <- predict(multinom_model, newdata = df_filtered[test,], type = "class")
confusion_matrix <- table(predictions, y.test)
accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
print(paste("Accuracy:", accuracy))

```

```
## [1] "Accuracy: 0.495555555555556"
```

We get an accuracy of around 50% relatively good improvement compared to the value of the 11.1% accuracy of a random guess.

```
# Display results
print("Confusion Matrix:")
```

```
## [1] "Confusion Matrix:"
```

```
print(confusion_matrix)
```

```
##           y.test
## predictions blues classical country electronic hip-hop jazz pop reggae rock
## blues       53         9      13         7        6   9  13      12   15
## classical    5       177       2         3        0  17   0       0   0
## country     11        0      89        11       4  20  13       7  46
## electronic  12        0       4        95       14  12  10       5  10
## hip-hop      3        0       3        32       81   1  34      46   2
## jazz        38       16      23         3        0 115  12       0  12
## pop         26        3      11        19       39   8  76      20  19
## reggae      14        0       7         9       54   5  26     131  12
## rock        27        2      45        15        4  10  14       4  75
```

Looking at the Confusion Matrix we and the overall precision values in the table below , we can notice that the classical genre is classified with a relatively high precision of 86.7% while the rest are all below 50% with the exception of electronic and jazz with 58.6% and 52.8% respectively. We can notice that The Blues genre is most confused by jazz, hip-hop is most confused by Reggae and rock is most confused by country.

```
# Print or use the metrics_df data frame
print(metrics_df)
```

```
##           Precision    Recall  F1_Score Specificity
## Class: blues      0.3868613 0.2804233 0.3251534  0.7195767
## Class: classical  0.8676471 0.8550725 0.8613139  0.1449275
## Class: country    0.4427861 0.4517766 0.4472362  0.5482234
## Class: electronic 0.5864198 0.4896907 0.5337079  0.5103093
## Class: hip-hop    0.4009901 0.4009901 0.4009901  0.5990099
## Class: jazz       0.5251142 0.5837563 0.5528846  0.4162437
## Class: pop        0.3438914 0.3838384 0.3627685  0.6161616
## Class: reggae     0.5077519 0.5822222 0.5424431  0.4177778
## Class: rock       0.3826531 0.3926702 0.3875969  0.6073298
```

3.1.2 Subset Selection

In this subsection, we delve into subset selection methods for model comparison, specifically focusing on backward subset selection using AIC, forward subset selection using AIC, backward selection using BIC, and forward selection using BIC. Subset selection methods are a class of techniques used to identify the most informative subset of predictor variables for a given statistical model. These methods aim to strike a balance between model complexity and predictive performance, ultimately leading to models that are more interpretable and generalize better to unseen data. The Akaike Information Criterion (AIC) and Bayesian

Information Criterion (BIC) are widely used model selection criteria that balance goodness of fit with model complexity. They are defined as follows:

$$\text{AIC} = -2\log(L) + 2k$$

$$\text{BIC} = -2\log(L) + k\log(n)$$

where L is the likelihood of the model, k is the number of parameters in the model, and n is the number of observations. The AIC and BIC penalize models with higher complexity, encouraging the selection of simpler models that adequately explain the data.

- **Backward selection with AIC**

Backward subset selection is a stepwise approach that starts with the full model and iteratively removes predictors that contribute the least to the model's fit. In backward subset selection with AIC, predictors are sequentially removed until the AIC is minimized, indicating the best subset of predictors.

```
model <- multinom(track_genre ~ ., data = df_filtered[train,])  
  
# Perform stepwise selection using step AIC function  
step_model <- step(model, direction = "backward")
```

```
step_model$anova
```

```
##      Step Df Deviance Resid. Df Resid. Dev      AIC  
## 1      NA      NA      7072    20199.25 20455.25  
## 2 - key   8 13.37702      7080    20212.63 20452.63
```

Only the key feature was removed. Testing the new model trained on the new model trained on a subset removing the key feature, the overall accuracy on a test set actually slightly decreased by around 1%.

```
predictions <- predict(step_model, newdata = df_filtered[test,], type = "class")  
confusion_matrix <- table(predictions, y.test)  
accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)  
print(paste("Accuracy:", accuracy))
```

```
## [1] "Accuracy: 0.4888888888888889"
```

From the confusion Matrix and the table of metrics, we can see that the main genre negatively affected by removing the key feature is the Blues genre with a reduction in precision of nearly 4% and its increase in confusion with jazz.

```
print("Confusion Matrix:")
```

```
## [1] "Confusion Matrix:"
```

```
print(confusion_matrix)
```

```
##           y.test
## predictions blues classical country electronic hip-hop jazz pop reggae rock
## blues      45         9      13         6         6      9 13      12 15
## classical   5        177      2         3         0     17  0         0  0
## country    10         0     86        11         4     19 13         7 45
## electronic 12         0      4        97        14     12 10         5 10
## hip-hop     3         0      3        31        81      1 33        47  2
## jazz       45        16     25         3         0    115 13         0 14
## pop        26         3     11        19        38      9 76        21 18
## reggae     14         0      6         8        55      5 26       129 13
## rock       29         2     47        16         4     10 14         4 74
```

```
# Print or use the metrics_df data frame
print(metrics_df)
```

```
##           Precision    Recall  F1_Score Specificity
## Class: blues      0.3515625 0.2380952 0.2839117 0.7619048
## Class: classical  0.8676471 0.8550725 0.8613139 0.1449275
## Class: country    0.4410256 0.4365482 0.4387755 0.5634518
## Class: electronic 0.5914634 0.5000000 0.5418994 0.5000000
## Class: hip-hop    0.4029851 0.4009901 0.4019851 0.5990099
## Class: jazz       0.4978355 0.5837563 0.5373832 0.4162437
## Class: pop        0.3438914 0.3838384 0.3627685 0.6161616
## Class: reggae     0.5039062 0.5733333 0.5363825 0.4266667
## Class: rock       0.3700000 0.3874346 0.3785166 0.6125654
```

- Forward selection with AIC

Forward subset selection is another stepwise approach that begins with an empty model and iteratively adds predictors that most improve the model fit. Similar to backward selection, forward selection utilizes the AIC criterion to determine the optimal subset of predictors by adding variables until the AIC is minimized.

```
output <- capture.output({
#Forward feature selection(AIC)
selected_features <- c()
min_AIC <- Inf

while (TRUE) {
  remaining_features <- setdiff(colnames(df_filtered), c("track_genre", selected_features))

  if (length(remaining_features) == 0) {
    break
  }

  best_AIC <- Inf
  best_feature <- NULL

  for (feature in remaining_features) {
    current_model <- multinom(track_genre ~ ., data = df_filtered[train, c(selected_features, feature,
    current_AIC <- AIC(current_model)

    if (current_AIC < best_AIC) {
      best_AIC <- current_AIC
```

```

    best_feature <- feature
  }
}

if (best_AIC < min_AIC) {
  min_AIC <- best_AIC
  selected_features <- c(selected_features, best_feature)
} else {
  break # No improvement in AIC, stop the loop
}
}

})

```

The same subset was obtained with forward selection, only key was removed.

```
print(selected_features)
```

```
## [1] "acousticness"      "danceability"      "instrumentalness"  "popularity"
## [5] "speechiness"       "valence"           "loudness"         "mode"
## [9] "energy"            "duration_ms"       "explicit"          "tempo"
## [13] "time_signature"    "liveness"
```

- Backward selection with BIC

Backward subset selection with the Bayesian Information Criterion (BIC) follows a similar approach to AIC-based selection but utilizes the BIC criterion instead. BIC penalizes model complexity more severely than AIC, often resulting in more parsimonious models. In backward subset selection with BIC, predictors are eliminated iteratively until the BIC is minimized, leading to a model with optimal predictive performance.

```
model <- multinom(track_genre ~ ., data = df_filtered[train,])

# Perform stepwise selection using step function
step_model <- step(model, direction = "backward", k=log(length(train)), trace=0)
```

```
step_model$anova
```

```
##      Step Df Deviance Resid. Df Resid. Dev      AIC
## 1      NA    NA      7072    20199.25 21336.13
## 2 - key   8 13.37702      7080    20212.63 21278.45
```

```
step_model$coefnames
```

```
## [1] "(Intercept)"      "popularity"        "duration_ms"       "explicitTrue"
## [5] "danceability"     "energy"            "loudness"          "mode"
## [9] "speechiness"      "acousticness"      "instrumentalness"  "liveness"
## [13] "valence"          "tempo"             "time_signature"
```

The same subset was selected, with only the “key” feature removed.

- Forward selection with BIC

Forward subset selection with BIC starts with an empty model and adds predictors based on the BIC criterion. Similar to the forward selection with AIC, this method iteratively adds predictors until the BIC is minimized, resulting in a subset of predictors that optimally balances model fit and complexity.

```
output<-capture.output({
selected_features <- c()
min_BIC <- Inf

while (TRUE) {
  remaining_features <- setdiff(colnames(df_filtered), c("track_genre", selected_features))

  if (length(remaining_features) == 0) {
    break
  }

  best_BIC <- Inf
  best_feature <- NULL

  for (feature in remaining_features) {
    current_model <- multinom(track_genre ~ ., data = df_filtered[train, c(selected_features, feature,
    current_BIC <- BIC(current_model)

    if (current_BIC < best_BIC) {
      best_BIC <- current_BIC
      best_feature <- feature
    }
  }

  if (best_BIC < min_BIC) {
    min_BIC <- best_BIC
    selected_features <- c(selected_features, best_feature)
  } else {
    break # No improvement in BIC, stop the loop
  }
}
})
```

```
length(selected_features)
```

```
## [1] 12
```

```
print(selected_features)
```

```
## [1] "acousticness"      "danceability"      "instrumentalness" "popularity"
## [5] "speechiness"      "valence"           "loudness"         "mode"
## [9] "energy"           "duration_ms"       "explicit"          "tempo"
```

With forward subset selection with BIC , 3 features were removed: key,liveness and mode. We can get a preliminary idea of the relative unimportance of these 3 features in the multiclass classification

```

multinom_model_bic <- multinom(track_genre ~ ., data = df_filtered[train,c(selected_features,"track_genre")])

# Make Predictions
predictions <- predict(multinom_model_bic, newdata = df_filtered[test,c(selected_features,"track_genre")])

# Evaluate the Model
confusion_matrix <- table(predictions, y.test)
accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)

print(paste("Accuracy:", accuracy))

```

```
## [1] "Accuracy: 0.493888888888889"
```

The accuracy on the model applied on the test set is slightly better than the previous subset selection of only removing key, but is still slightly less by an insignificant 0.5% compared to the full model.

```
print("Confusion Matrix:")
```

```
## [1] "Confusion Matrix:"
```

```
print(confusion_matrix)
```

```
##           y.test
## predictions blues classical country electronic hip-hop jazz pop reggae rock
## blues       53         10        14           6          8    9  16    13   12
## classical    5        178         2           3          0   17   0     0    0
## country     12         0        96          12         4   22  13     5   46
## electronic  11         0         4          93        12  12  10     5   11
## hip-hop      3         0         3          31        80   1  32    47    2
## jazz        41        15        26           2          0  113  12     0   14
## pop         23         2        12          22        37   9  71    19   18
## reggae      15         0         7           7        58   5  27   130   13
## rock        26         2        33          18         3   9  17     6   75
```

```

# Print or use the metrics_df data frame
print(metrics_df)

```

```
##           Precision    Recall  F1_Score Specificity
## Class: blues      0.3758865 0.2804233 0.3212121  0.7195767
## Class: classical  0.8682927 0.8599034 0.8640777  0.1400966
## Class: country    0.4571429 0.4873096 0.4717445  0.5126904
## Class: electronic 0.5886076 0.4793814 0.5284091  0.5206186
## Class: hip-hop    0.4020101 0.3960396 0.3990025  0.6039604
## Class: jazz       0.5067265 0.5736041 0.5380952  0.4263959
## Class: pop        0.3333333 0.3585859 0.3454988  0.6414141
## Class: reggae     0.4961832 0.5777778 0.5338809  0.4222222
## Class: rock       0.3968254 0.3926702 0.3947368  0.6073298
```

Overall by rounding the accuracies of all 3 models seen till now, no difference was found in accuracy, while it looks like mode, key and liveness features have an insignificant effect on the multiclass classification.

3.1.3 Ridge Regularization

Ridge regression is a type of linear regression that introduces regularization to mitigate multicollinearity and reduce the variance of the parameter estimates. It adds a penalty term to the least squares objective function, constraining the magnitude of the coefficient estimates. In this subsection, we apply ridge regularization using the `glmnet` package in R, setting the tuning parameter λ to control the strength of regularization.

Ridge regression minimizes the following objective function:

$$\text{minimize } (\|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda\|\beta\|_2^2)$$

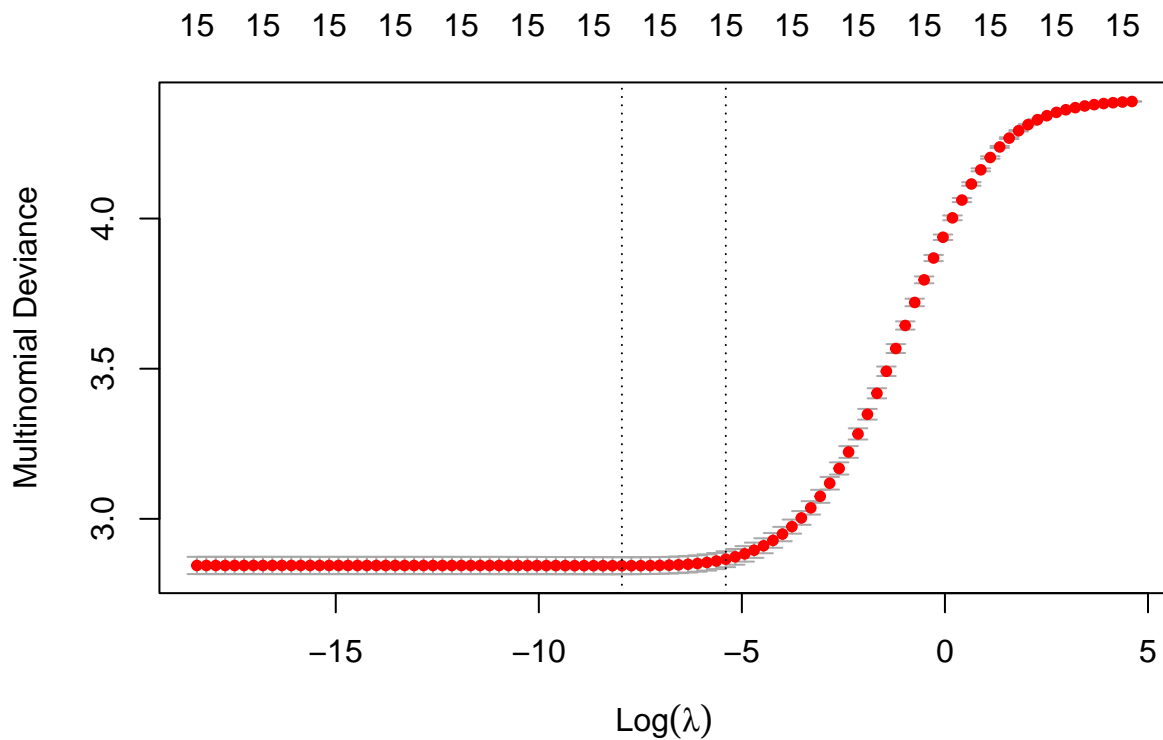
where $\|\mathbf{y} - \mathbf{X}\beta\|_2^2$ is the residual sum of squares, $\|\beta\|_2^2$ is the L2-norm penalty term, and λ is the regularization parameter. The L2-norm penalty term shrinks the coefficient estimates towards zero, reducing their variance and mitigating overfitting.

The `glmnet` package provides efficient algorithms for fitting generalized linear models with elastic-net regularization, which combines ridge (L2) and lasso (L1) penalties. However, by setting the parameter `alpha` to 0, we exclusively apply ridge regularization in our model.

Through ridge regularization, we aim to stabilize the parameter estimates and improve the generalization performance of our model. We evaluate the effectiveness of ridge regularization by tuning the regularization parameter λ by cross-validation. By incorporating ridge regularization into our modeling pipeline, we strive to strike a balance between model complexity and predictive accuracy.

```
fit_ridge<-glmnet(X_scaled[train,],y[train],family="multinomial",alpha=0)
#plot(fit_ridge)
```

```
lambda_interval<-10^seq(2,-8,length=100)
cvfit_ridge<-cv.glmnet(X_scaled[train,],y[train],family="multinomial",alpha=0,lambda=lambda_interval)
plot(cvfit_ridge)
```



```
#No Regularization
pred <- predict(fit_ridge, s = 0, newx = X_scaled[test, ],
                exact = TRUE, type="class", x = X_scaled[train, ], y = y[train])
pred_accuracy <- mean(pred == y.test)
pred_accuracy
```

```
## [1] 0.4938889
```

```
bestlam <- cvfit_ridge$lambda.min
bestlam
```

```
## [1] 0.0003511192
```

```
ridge_pred = predict(cvfit_ridge, newx = X_scaled[test, ], type = "class", s = "lambda.min")
ridge_pred_accuracy <- mean(ridge_pred == y.test)
ridge_pred_accuracy
```

```
## [1] 0.4944444
```

```
conf_matrix <- table(Actual = y.test, Predicted = ridge_pred)
print(conf_matrix)
```

```
##          Predicted
```

## Actual	blues	classical	country	electronic	hip-hop	jazz	pop	reggae	rock
## blues	51	5	10	12	4	38	26	14	29
## classical	9	177	0	0	0	16	3	0	2
## country	13	1	88	4	3	25	11	7	45
## electronic	7	3	11	95	32	3	19	9	15
## hip-hop	6	0	4	15	82	0	36	55	4
## jazz	7	17	21	12	1	116	8	5	10
## pop	12	0	13	10	33	12	77	27	14
## reggae	13	0	7	6	47	0	19	129	4
## rock	14	0	44	10	2	14	19	13	75

```
# Print or use the metrics_df data frame
print(metrics_df)
```

##	Precision	Recall	F1_Score	Specificity
## Class: blues	0.2698413	0.3863636	0.3177570	0.6136364
## Class: classical	0.8550725	0.8719212	0.8634146	0.1280788
## Class: country	0.4467005	0.4444444	0.4455696	0.5555556
## Class: electronic	0.4896907	0.5792683	0.5307263	0.4207317
## Class: hip-hop	0.4059406	0.4019608	0.4039409	0.5980392
## Class: jazz	0.5888325	0.5178571	0.5510689	0.4821429
## Class: pop	0.3888889	0.3532110	0.3701923	0.6467890
## Class: reggae	0.5733333	0.4980695	0.5330579	0.5019305
## Class: rock	0.3926702	0.3787879	0.3856041	0.6212121

3.1.4 Lasso Regularization

Lasso (Least Absolute Shrinkage and Selection Operator) regression is another form of linear regression that introduces regularization to select a subset of relevant predictors and shrink the coefficients of less important predictors to zero. This technique helps in building parsimonious models and handling multicollinearity.

Lasso regression minimizes the following objective function:

$$\text{minimize } (\|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda\|\beta\|_1)$$

where $\|\mathbf{y} - \mathbf{X}\beta\|_2^2$ is the residual sum of squares, $\|\beta\|_1$ is the L1-norm penalty term, and λ is the regularization parameter.

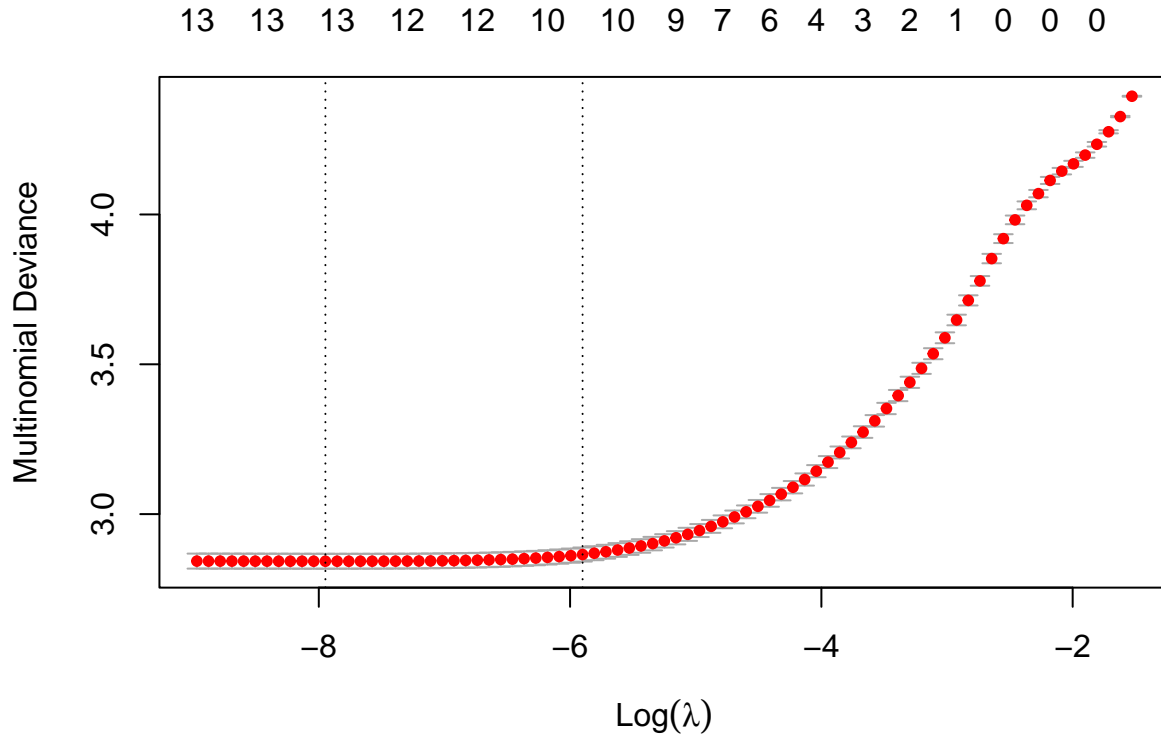
In this subsection, we apply Lasso regularization using the `glmnet` package in R. By setting the parameter `alpha` to 1, we exclusively apply Lasso regularization in our model. The L1-norm penalty encourages sparsity in the coefficient estimates, effectively performing variable selection by driving some coefficients to zero.

The tuning parameter λ controls the strength of regularization and the degree of shrinkage applied to the coefficient estimates. Through cross-validation we select the optimal value of λ that balances model complexity and predictive performance.

We evaluate the effectiveness of Lasso regularization by examining the resulting coefficient estimates, identifying significant predictors, and assessing the model's performance on validation data. Through this approach, we strive to build robust and parsimonious models that capture the underlying patterns in the data while minimizing overfitting.

```
fit_lasso<-glmnet(X_scaled[train,],y[train],family="multinomial",alpha=1)
custom_lambda_values <- 10^seq(10, -4, length = 20) # Customize the range and number of lambda values
```

```
cvfit_lasso <- cv.glmnet(X_scaled[train,], y[train], family = "multinomial", alpha = 1)
plot(cvfit_lasso)
```



```
bestlam.1se <- cvfit_lasso$lambda.1se
bestlam.1se
```

```
## [1] 0.002737711
```

```
lasso_coefficients<-coef(fit_lasso,s=bestlam.1se)
non_zero_counts <- sapply(lasso_coefficients, function(mat) sum(nnzero(mat)))
```

```
# Print the non-zero counts
print(non_zero_counts)
```

```
##      blues  classical  country electronic  hip-hop  jazz  pop
##       13       8       9       12       11       14      13
##    reggae      rock
##       11       9
```

```
lasso_1se_pred=predict(cvfit_lasso,newx=X_scaled[test,],type="class", s="lambda.1se")
lasso_1se_pred_accuracy<-mean(lasso_1se_pred==y.test)
lasso_1se_pred_accuracy
```

```
## [1] 0.4938889
```

```
conf_matrix <- table(Actual = y.test, Predicted = lasso_1se_pred)
print(conf_matrix)
```

```
##               Predicted
## Actual      blues classical country electronic hip-hop jazz pop reggae rock
## blues        49         4      10         10      4  41  26   15  30
## classical    10        175      0         0      0  17   3    0   2
## country      12         1      87         4      3  28  13    7  42
## electronic   7         4      10        91     34   2  20   12  14
## hip-hop      7         0       5        15     83   0  34   54   4
## jazz         6        17     19        12      1 122   7    6   7
## pop          11         0     12        13     31  12  74   27  18
## reggae       11         0      3         4     46   0  22  134   5
## rock         13         0     43        13      1  17  19   11  74
```

```
# Print or use the metrics_df data frame
print(metrics_df)
```

```
##               Class Precision   Recall  F1_Score Specificity
## Class: blues      Class: blues 0.2592593 0.3888889 0.3111111 0.6111111
## Class: classical  Class: classical 0.8454106 0.8706468 0.8578431 0.1293532
## Class: country    Class: country 0.4416244 0.4603175 0.4507772 0.5396825
## Class: electronic Class: electronic 0.4690722 0.5617284 0.5112360 0.4382716
## Class: hip-hop    Class: hip-hop 0.4108911 0.4088670 0.4098765 0.5911330
## Class: jazz       Class: jazz 0.6192893 0.5104603 0.5596330 0.4895397
## Class: pop        Class: pop 0.3737374 0.3394495 0.3557692 0.6605505
## Class: reggae     Class: reggae 0.5955556 0.5037594 0.5458248 0.4962406
## Class: rock       Class: rock 0.3874346 0.3775510 0.3824289 0.6224490
```

3.1.4.1 Lasso Coefficients In this subsection, we employ a bootstrap approach to assess the stability and variability of Lasso coefficients across different samples. Bootstrap resampling involves randomly sampling the dataset with replacement to generate multiple bootstrap samples. We repeat this process for a specified number of iterations, 1000 bootstraps in this case, to obtain a distribution of parameter estimates.

We apply the Lasso regression using the `glmnet` package in R to each bootstrap sample and extract the coefficients for each predictor variable. By aggregating the coefficient estimates across all bootstrap samples, we obtain insights into the robustness and significance of predictor variables in the Lasso model.

To visualize the distribution of Lasso coefficients for each class, we create box plots where the y-axis represents the coefficient values and the x-axis represents the predictor variables. Each box plot corresponds to a class, providing insights into the magnitude and variability of coefficients across different classes.

Through this analysis, we aim to identify important predictors that consistently contribute to the classification of each class while accounting for the variability introduced by the bootstrap resampling process. Box plots serve as effective visualizations to compare coefficient distributions and detect potential outliers or influential predictors.

```
library(doParallel)
library(foreach)
num_cores <- detectCores()
cl <- makeCluster(num_cores)
registerDoParallel(cl)
```

```

# Define the number of bootstraps
num_bootstraps <- 1000

# Define the function to get coefficients from bootstrapped samples
get_coefs <- function( indices) {
  #coef(model, s = cv.glmnet(X_scaled.train[indices, ], y.train[indices], family = "multinomial", alpha
  library(glmnet)
  model <- glmnet(
    x = X_scaled.train[indices, ],
    y = y.train[indices],
    family = "multinomial",
    alpha=1,
    parallel = TRUE
  )

  # Extract and return the coefficients as a matrix
  #return(as.matrix(coef(model,s=bestlam.1se)))
  coefs_list <- coef(model, s = bestlam.1se)

  # Convert the list of coefficient vectors to a list of matrices
  coefs_matrices <- lapply(coefs_list, as.matrix)

  return(coefs_matrices)
}

# Bootstrap procedure in parallel
boot_coefs <- foreach(i = 1:num_bootstraps, .combine = "rbind") %dopar% {
  set.seed(i) # Set a seed for reproducibility
  indices <- sample(1:nrow(X_scaled.train), nrow(X_scaled.train), replace = TRUE)
  get_coefs( indices)
}

# Stop the parallel backend
stopCluster(cl)

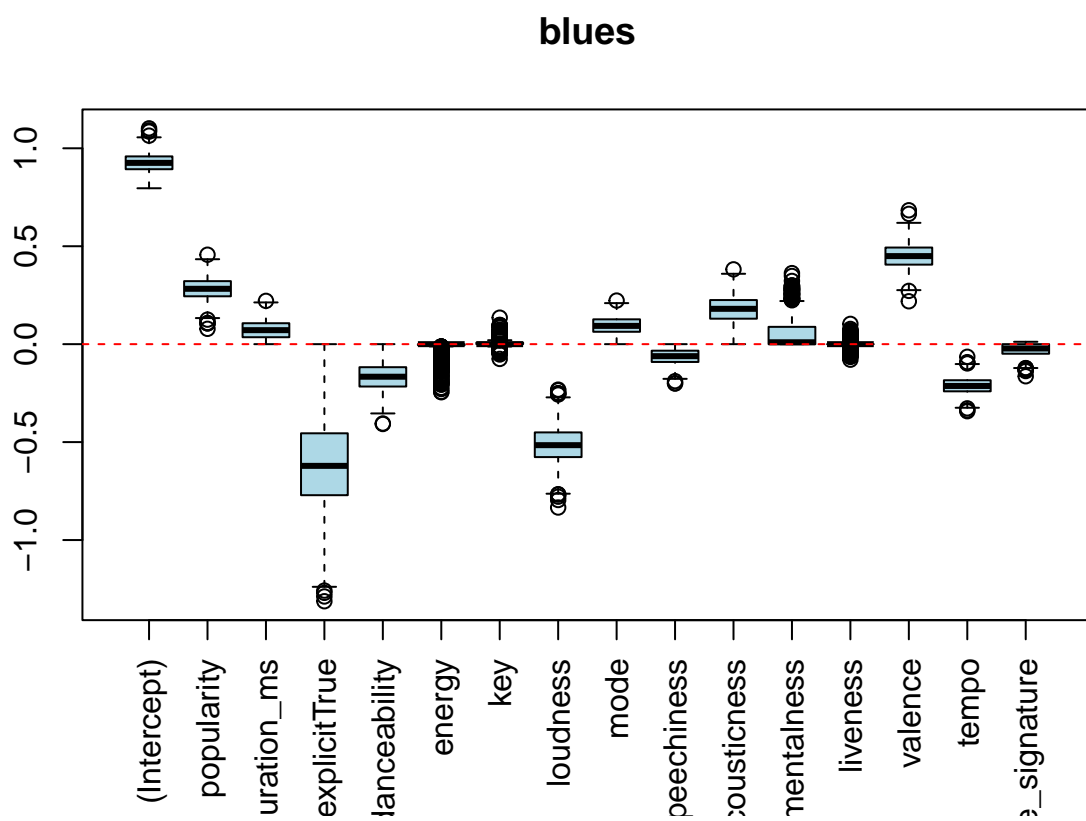
```

1. Blues

```

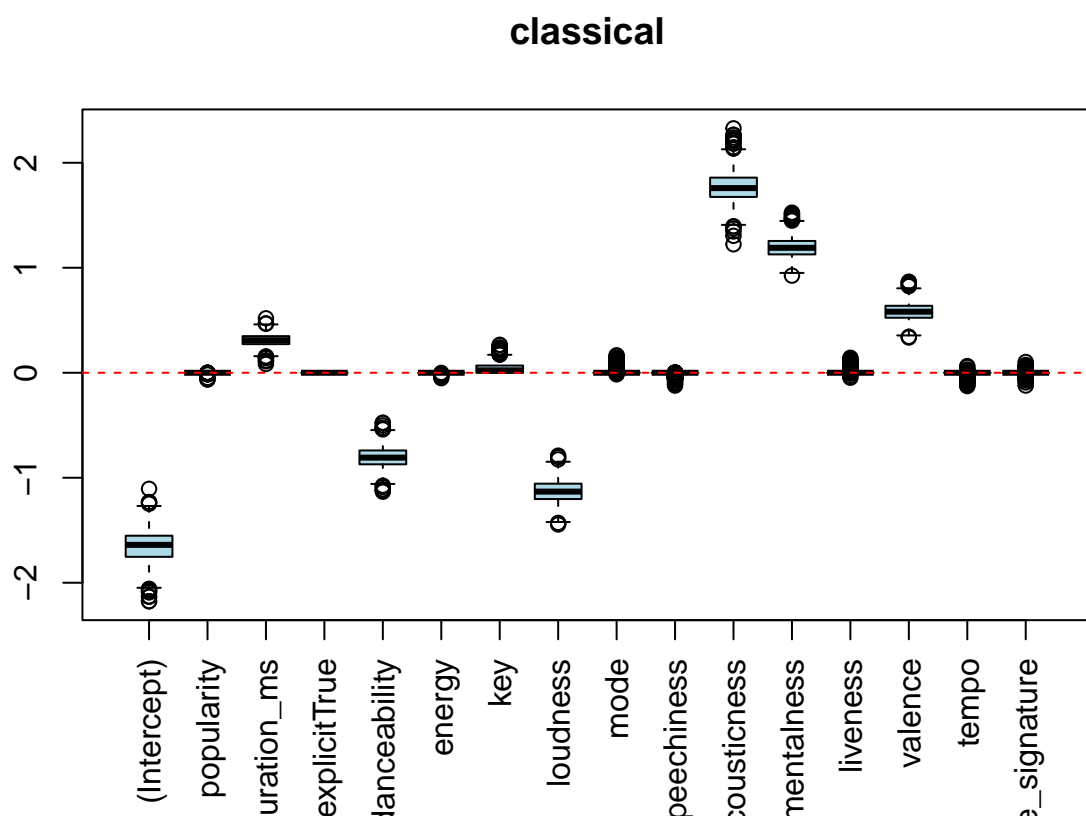
#par(mfrow = c(3,3))
i<-1
#for (i in (1:ncol(boot_coefs))) {
  class_coefs <- boot_coefs[,i]
  class_coefs_c<-do.call(cbind,class_coefs)
  boxplot(t(class_coefs_c), main = colnames(boot_coefs)[i],
    col = "lightblue",
    names = colnames(class_coefs),xaxt="n")
  abline(h=0,col="red",lty=2)
  axis(1, at = seq_along(rownames(class_coefs_c)), labels = rownames(class_coefs_c), las = 2)
}

```

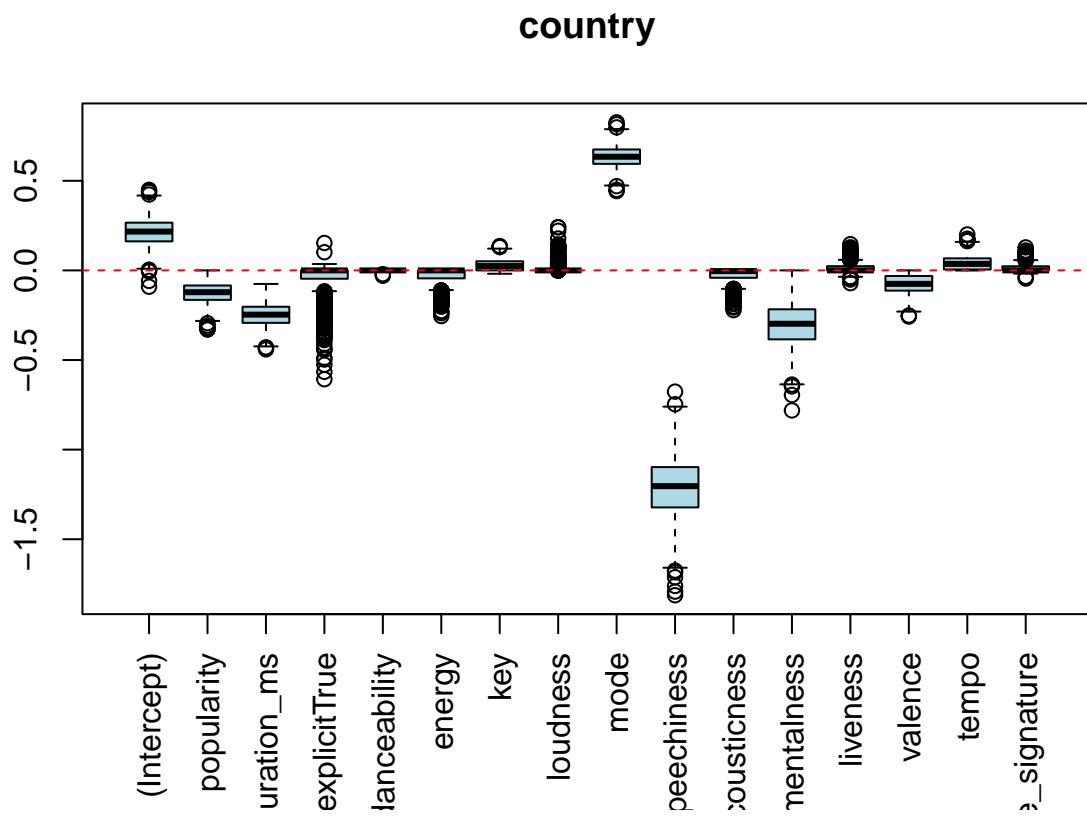



```
#}
#par(mfrow=c(1,1))
```

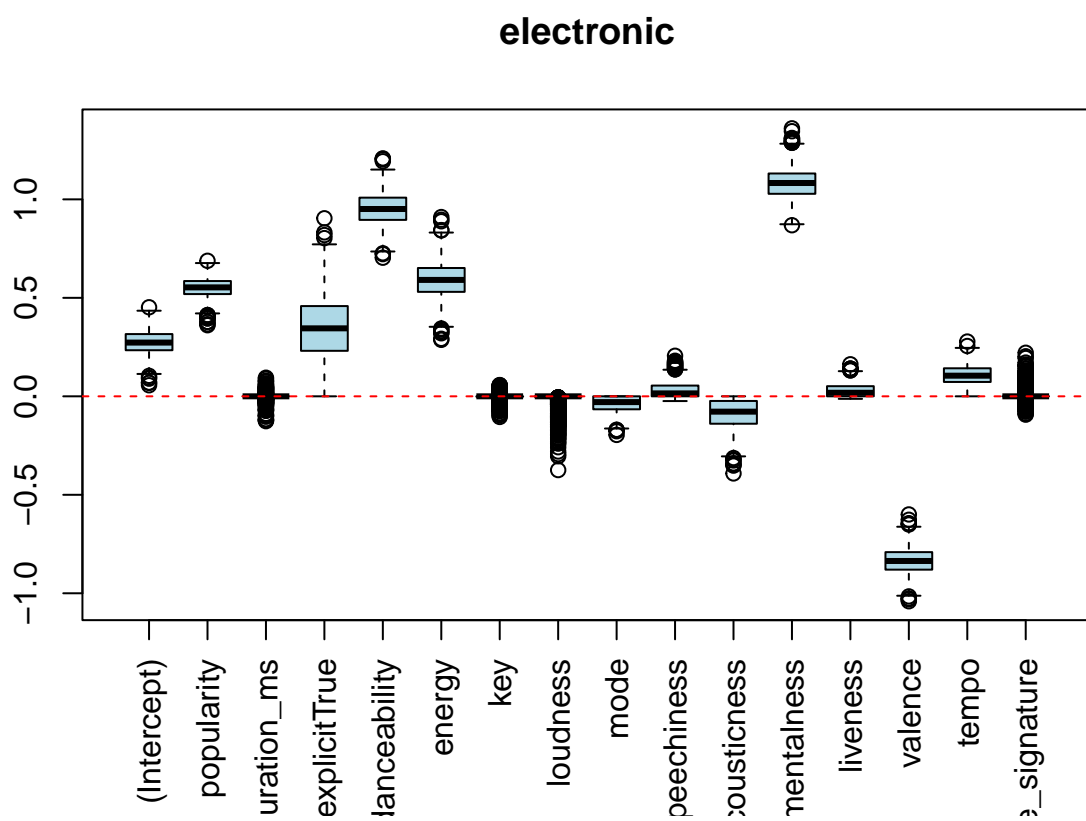
2. Classical



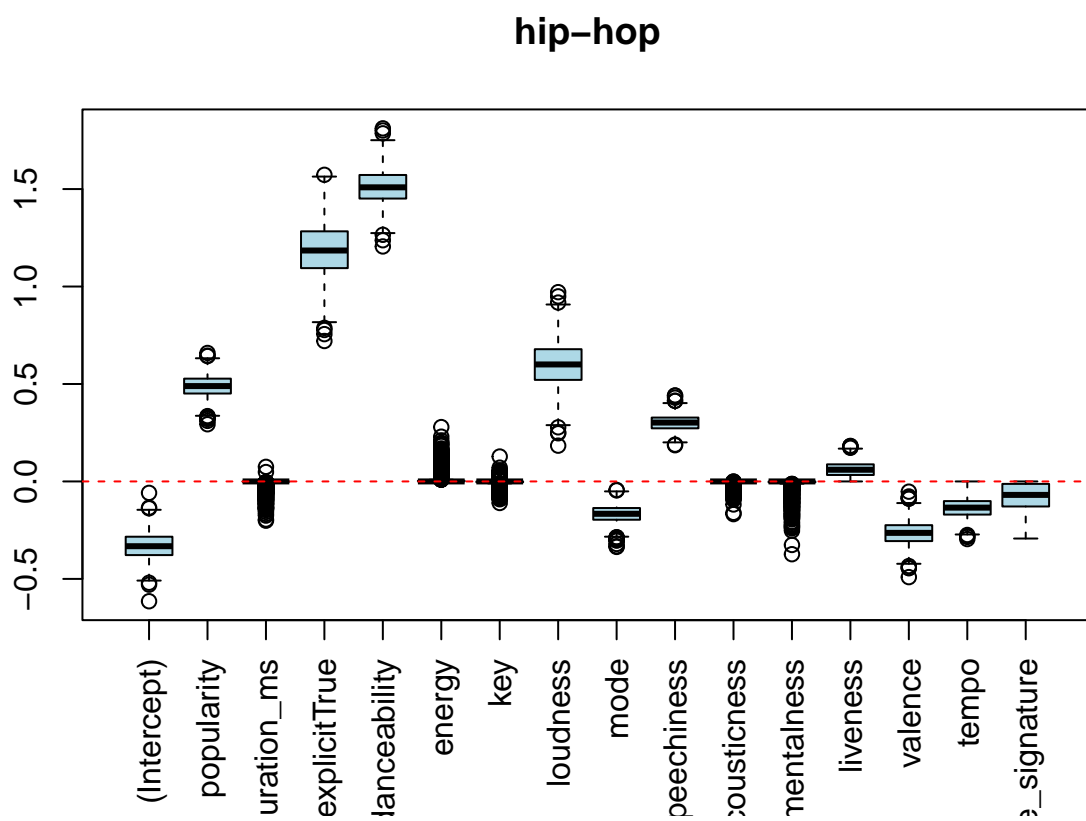
3. Country



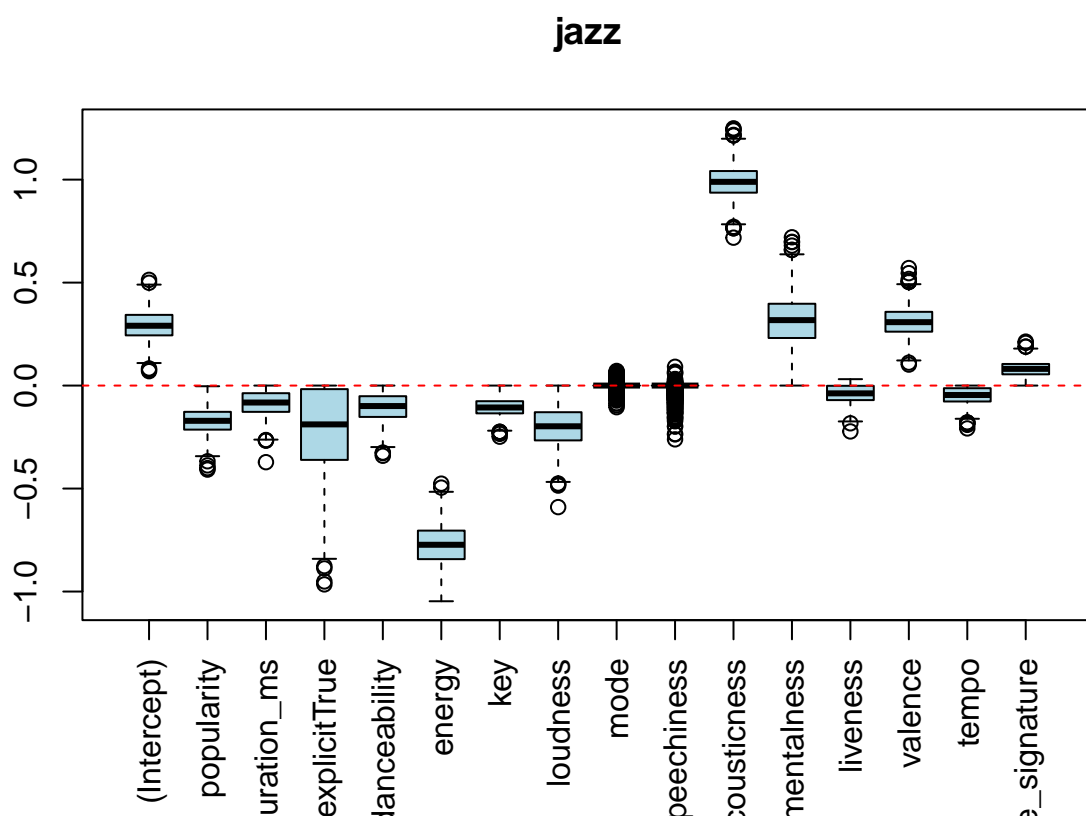
4. Electronic



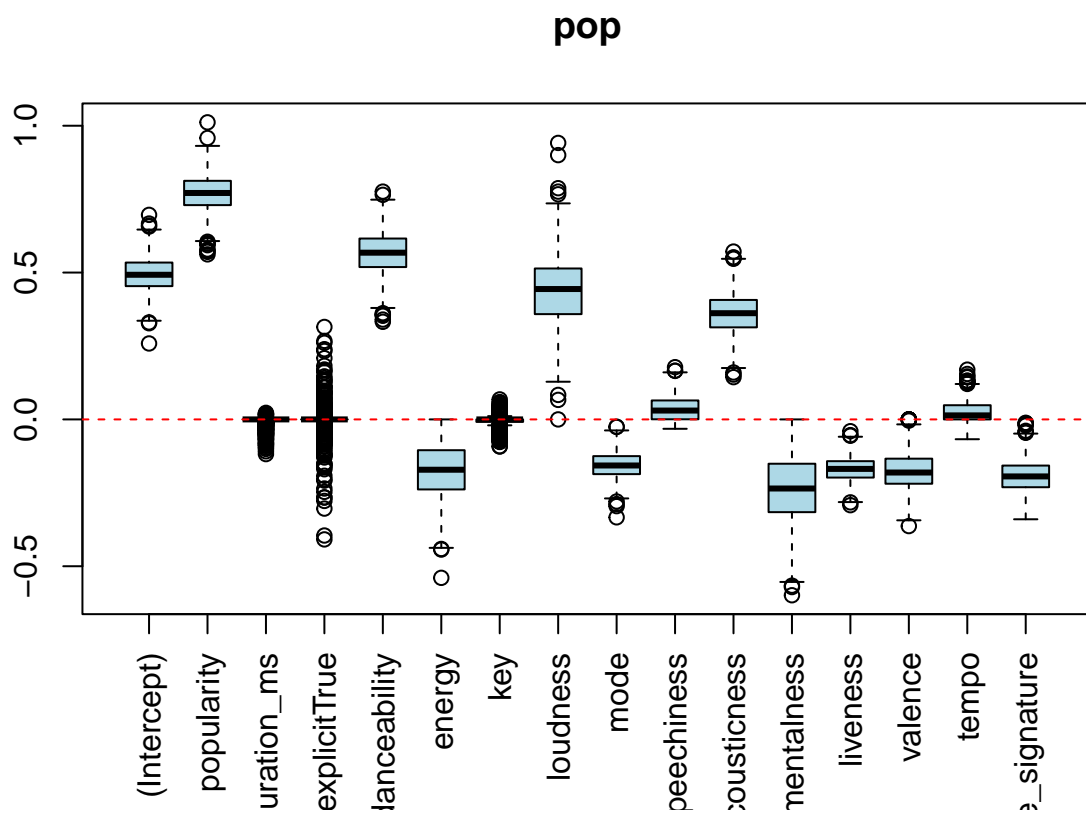
5. Hip-Hop



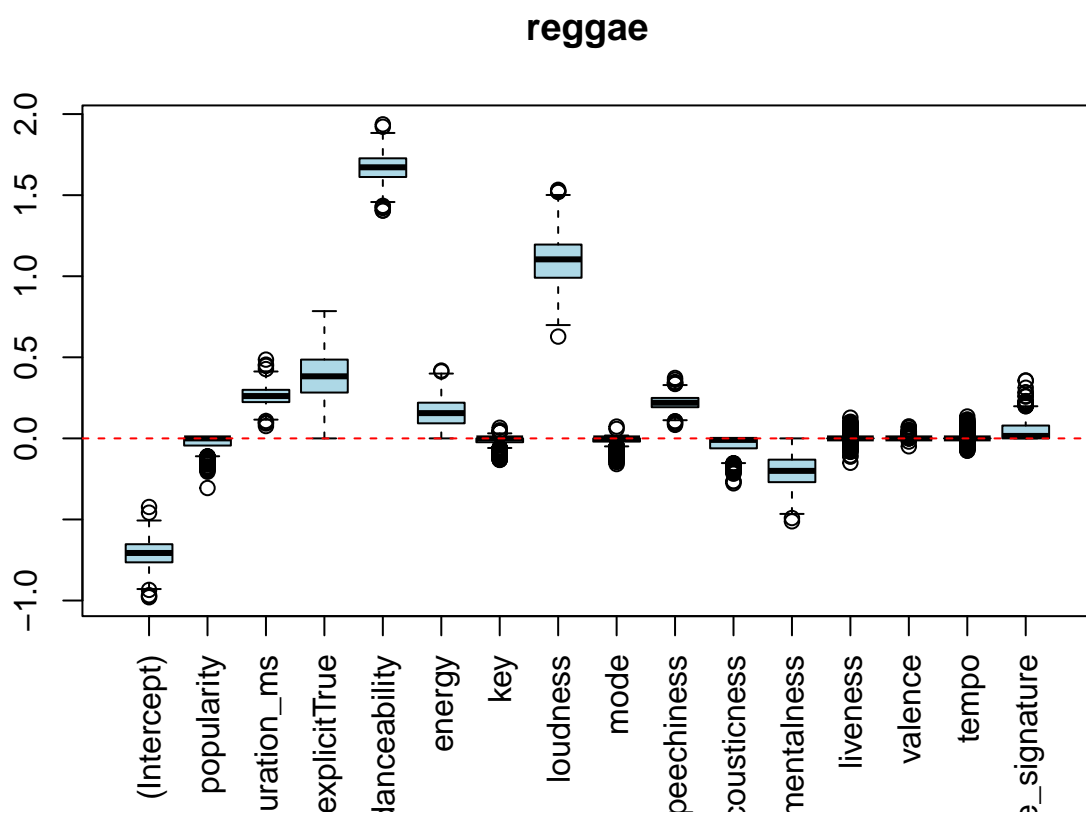
6. Jazz



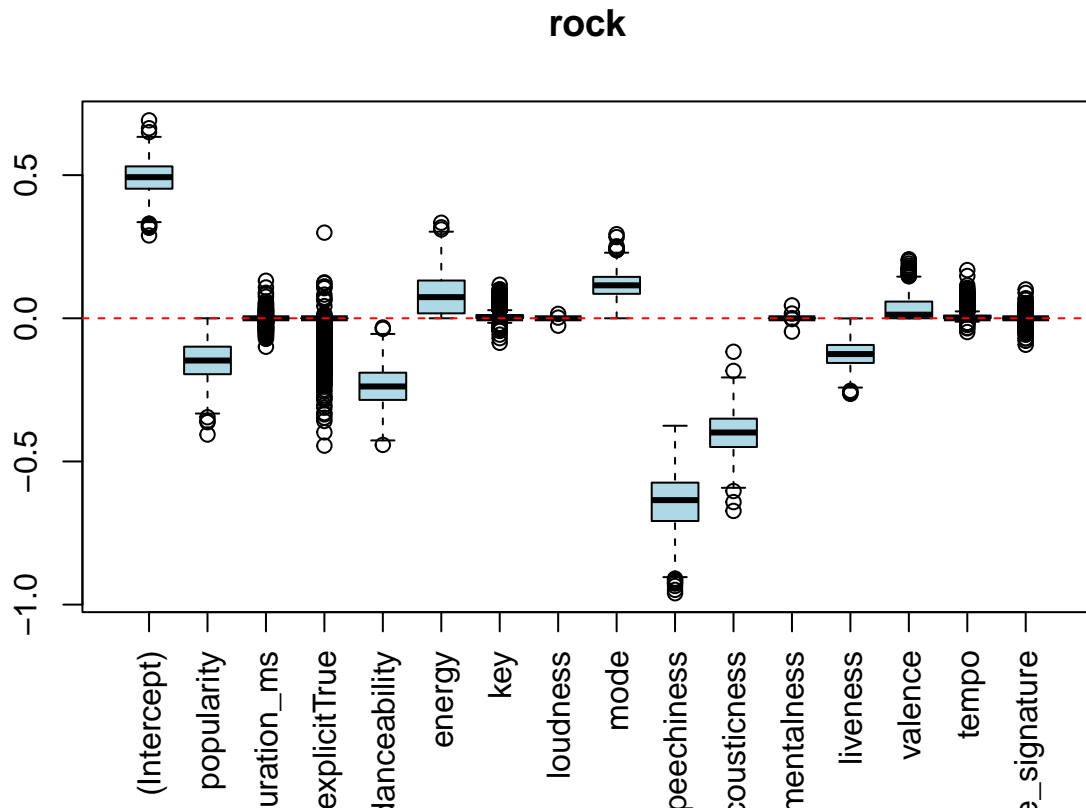
7. Pop



8. Reggae



9. Rock



3.1.5 Elastic Net

Elastic Net regularization is a hybrid of ridge and Lasso regularization techniques, combining the strengths of both methods to handle multicollinearity, perform variable selection, and shrink coefficient estimates. Elastic Net introduces two tuning parameters, α and λ , controlling the balance between the L1-norm and L2-norm penalties.

The objective function for Elastic Net regularization is defined as:

$$\text{minimize } (\|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda (\alpha \|\beta\|_1 + (1 - \alpha) \|\beta\|_2^2))$$

where $\|\mathbf{y} - \mathbf{X}\beta\|_2^2$ is the residual sum of squares, $\|\beta\|_1$ is the L1-norm penalty term, $\|\beta\|_2^2$ is the L2-norm penalty term, λ is the regularization parameter, and α controls the mixing ratio between L1 and L2 penalties.

In this subsection, we apply Elastic Net regularization using the `glmnet` package in R. By setting α to a value between 0 and 1, we adjust the balance between Lasso and ridge penalties. A value of $\alpha = 1$ corresponds to pure Lasso regularization, while $\alpha = 0$ corresponds to pure ridge regularization.

Through cross-validation or other validation techniques, we select the optimal values of α and λ that balance model complexity and predictive performance. Elastic Net regularization is particularly useful when dealing with high-dimensional data and highly correlated predictor variables. Through this approach, we strive to harness the advantages of both ridge and Lasso regularization techniques.

```
set.seed(1)
# Initialize empty lists
alpha_list <- numeric()
```

```

error_min_list <- numeric()
error_1se_list <- numeric()
lambda_min_list <- numeric()
lambda_1se_list <- numeric()

#custom_lambda_sequence <- 10^seq(1, -6, length = 100)
for (a in seq(0, 1, by = 0.1)) {

  cvfit_elastic_net <- cv.glmnet(X_scaled[train,],y[train],family="multinomial", alpha = a)#,lambda=cus
  #plot(cvfit_elastic_net)
  #title(main = bquote(alpha == .(a)))
  # Extract alpha, lambda.min, lambda.1se, and cross-validated errors
  alpha_value <- cvfit_elastic_net$glmnet.fit$alpha
  lambda_min <- cvfit_elastic_net$lambda.min
  lambda_1se <- cvfit_elastic_net$lambda.1se
  error_min <- cvfit_elastic_net$cvm[cvfit_elastic_net$lambda == lambda_min]
  error_1se <- cvfit_elastic_net$cvm[cvfit_elastic_net$lambda == lambda_1se]

  # Append values to lists
  alpha_list <- c(alpha_list, a)
  error_min_list <- c(error_min_list, error_min)
  error_1se_list <- c(error_1se_list, error_1se)
  lambda_min_list <- c(lambda_min_list, lambda_min)
  lambda_1se_list <- c(lambda_1se_list, lambda_1se)
}

results_df <- data.frame(
  alpha = seq(0, 1, by = 0.1),
  error_min = error_min_list,
  error_1se = error_1se_list,
  lambda_min = lambda_min_list,
  lambda_1se = lambda_1se_list
)

print(results_df)

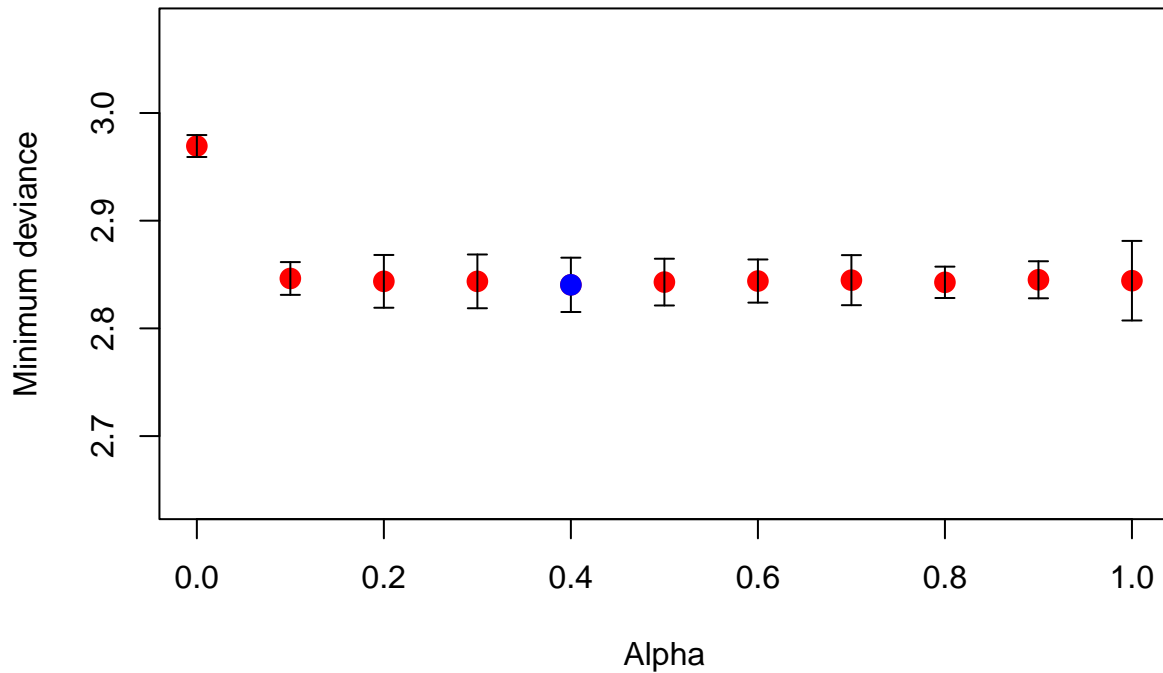
```

```

##      alpha error_min error_1se  lambda_min  lambda_1se
## 1      0.0  2.969325  2.979513 0.0216958893 0.023811212
## 2      0.1  2.846341  2.861553 0.0004566777 0.003535889
## 3      0.2  2.843672  2.868176 0.0003990286 0.004482384
## 4      0.3  2.843657  2.868612 0.0004648759 0.004335448
## 5      0.4  2.840427  2.865648 0.0003486569 0.003916547
## 6      0.5  2.842974  2.864686 0.0004441283 0.003438724
## 7      0.6  2.843949  2.863980 0.0005369618 0.003144996
## 8      0.7  2.844777  2.868012 0.0005051271 0.003246994
## 9      0.8  2.842763  2.857330 0.0004419862 0.002358747
## 10     0.9  2.845128  2.862313 0.0004311816 0.002525439
## 11     1.0  2.844283  2.881225 0.0005129970 0.003619092

```

Min binomial deviance vs. Alpha



```
alpha_chosen<-results_df$alpha[which.min(results_df$error_min)]

best_lam_elastic_net<-results_df$lambda_1se[which.min(results_df$error_min)]

best_elastic_net<-glmnet(X_scaled[train,],y[train],family="multinomial",alpha=alpha_chosen)

elastic_net_pred<-predict(best_elastic_net,newx=X_scaled.test,type="class",s=best_lam_elastic_net)
elastic_net_pred_accuracy<-mean(elastic_net_pred==y.test)
elastic_net_pred_accuracy
```

```
## [1] 0.4955556
```

```
conf_matrix <- table(Actual = y.test, Predicted = elastic_net_pred)
print(conf_matrix)
```

```
##          Predicted
## Actual    blues classical country electronic hip-hop jazz pop reggae rock
## blues      51         4      10         12         4   39  25   14   30
## classical   9       175       0         0         0   18   3    0    2
## country    12         1      85         4         3   28  13    7   44
## electronic  9         4       9         90        35   1  20   11   15
## hip-hop     7         0       4         15        83   0  33   56    4
## jazz        7        17      20         12         1 123   6    5    6
## pop         11         0      11         13        31  12  78   26   16
```

```
## reggae      11      0      4      5      47      0      22      132      4
## rock        14      0     42     11      2      16      18       13     75
```

```
# Print or use the metrics_df data frame
print(metrics_df)
```

```
##              Precision    Recall  F1_Score Specificity
## Class: blues      0.2698413 0.3893130 0.3187500  0.6106870
## Class: classical  0.8454106 0.8706468 0.8578431  0.1293532
## Class: country    0.4314721 0.4594595 0.4450262  0.5405405
## Class: electronic 0.4639175 0.5555556 0.5056180  0.4444444
## Class: hip-hop    0.4108911 0.4029126 0.4068627  0.5970874
## Class: jazz       0.6243655 0.5189873 0.5668203  0.4810127
## Class: pop        0.3939394 0.3577982 0.3750000  0.6422018
## Class: reggae     0.5866667 0.5000000 0.5398773  0.5000000
## Class: rock       0.3926702 0.3826531 0.3875969  0.6173469
```

3.2 K-Nearest Neighbours

K-Nearest Neighbors (KNN) is a non-parametric classification algorithm used for both classification and regression tasks. In KNN, the class of an observation is determined by the majority class among its K nearest neighbors in the feature space. KNN is particularly useful when dealing with non-linear relationships and high-dimensional data, making it a versatile choice for various classification tasks.

In this section, we apply KNN using the `knn()` function from the “class” package in R. We use the optimal value of K selected through cross-validation to build the final KNN model. The `knn()` function assigns the class label of the majority of the K nearest neighbors to each observation in the test dataset.

To choose the optimal value of K in KNN, we employ cross-validation using the `trainControl()` function from the caret package.

```
data.train <- data.frame(X_scaled.train, y.train)

# Define the training control with 10-fold cross-validation
ctrl <- trainControl(method = "cv", number = 10)

# Specify the tuning grid for k (e.g., values from 1 to 10)
grid <- expand.grid(k = 1:10)

# Train the KNN model with cross-validation
set.seed(123) # Set a seed for reproducibility
knn_model <- train(y.train ~ ., data = data.train, method = "knn", trControl = ctrl, tuneGrid = grid)
```

The optimal K found from cross-validation is K=1

```
optimal_k <- knn_model$bestTune$k
knn_model <- knn(train = X_scaled[train,], test = X_scaled[test,], cl = df_filtered[train,16], k = optimal_k)

# Evaluate the model
conf_matrix <- table(Actual = y.test, Predicted = knn_model)
print(conf_matrix)
```

```
##               Predicted
## Actual      blues classical country electronic hip-hop jazz pop reggae rock
## blues        70         2      19         19       7   34   9       7   22
## classical     8        181       1         2       1   11   2       0   1
## country      22         2     115         4       0    5   9      10  30
## electronic   20         4      12        101      23   2  14      11   7
## hip-hop       6         0       3         18      82   2  50      36   5
## jazz         43        17       4          7       3  108   4       4   7
## pop          11         2      13          7      50   2  85      16  12
## reggae        9         0       4          8      35   2  34     120  13
## rock          9         1      26          5       5    5  18       5  117
```

```
Accuracy<-mean(knn_model==y.test)
Accuracy
```

```
## [1] 0.5438889
```

```
# Print or use the metrics_df data frame
print(metrics_df)
```

```
##               Precision    Recall  F1_Score Specificity
## Class: blues      0.3703704 0.3535354 0.3617571  0.6464646
## Class: classical  0.8743961 0.8660287 0.8701923  0.1339713
## Class: country    0.5837563 0.5837563 0.5837563  0.4162437
## Class: electronic 0.5206186 0.5906433 0.5534247  0.4093567
## Class: hip-hop    0.4059406 0.3980583 0.4019608  0.6019417
## Class: jazz       0.5482234 0.6315789 0.5869565  0.3684211
## Class: pop        0.4292929 0.3777778 0.4018913  0.6222222
## Class: reggae     0.5333333 0.5741627 0.5529954  0.4258373
## Class: rock       0.6125654 0.5467290 0.5777778  0.4532710
```

3.3 Random Forest

Random Forest is a powerful ensemble learning method that constructs a multitude of decision trees during training and outputs the mode of the classes (classification) or mean prediction (regression) of the individual trees. It's known for its robustness, scalability, and ability to handle high-dimensional data with complex interactions.

To optimize the performance of the Random Forest model and prevent overfitting, we utilize cross-validation techniques. We split the dataset into training and validation sets using the `trainControl()` function from the `caret` package. By varying parameters such as the number of trees and the number of variables randomly sampled as candidates at each split, we tune the model for optimal performance.

In this section, we apply Random Forest using the `randomForest()` function from the “randomForest” package in R. We specify the number of trees in the forest and other hyperparameters based on cross-validation results. The `randomForest()` function constructs the ensemble of decision trees and aggregates their predictions to make final classifications or predictions.

We evaluate the performance of the Random Forest model using appropriate evaluation metrics such as accuracy, precision, recall, and F1-score.

```

#library(randomForest)

#library(caret)

# Create a training control with cross-validation
ctrl <- trainControl(method = "cv", number = 10, classProbs = TRUE)
levels(df_filtered$track_genre) <- make.names(levels(df_filtered$track_genre))
# Train a Random Forest model using cross-validation
rf_model <- train(track_genre ~ ., data = df_filtered[train,],
                  method = "rf",
                  trControl = ctrl)

```

From cross-validation we set mtry=8. “mtry” refers to the number of variables randomly sampled as candidates at each split when building each tree in the forest.

```

predictions <- predict(rf_model, newdata = df_filtered[test,])

# Evaluate the model
confusion_matrix <- table(predictions, y.test)
confusion_matrix

```

```

##           y.test
## predictions  blues classical country electronic hip-hop jazz pop reggae rock
## blues        88         8        15         17         4    38    5         0    3
## classical     4       192         1         4         0    18    0         0    0
## country      11         0       128         2         1     4    3         6   23
## electronic   23         4         5       146         7    16    1         6    5
## hip.hop       6         0         2        10       129     2   31        30    2
## jazz         28         3         7         0         1   112     6         0    4
## pop          10         0         9         6        34     3  129        24   11
## reggae        5         0         5         5        24     2   16       157    4
## rock         14         0        25         4         2     2    7         2  139

```

```

accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
print(paste("Accuracy:", accuracy))

```

```
## [1] "Accuracy: 0.677777777777778"
```

```

# Print or use the metrics_df data frame
print(metrics_df)

```

```

##           Precision    Recall  F1_Score Specificity
## Class: blues      0.4943820 0.4656085 0.4795640 0.53439153
## Class: classical  0.8767123 0.9275362 0.9014085 0.07246377
## Class: country    0.7191011 0.6497462 0.6826667 0.35025381
## Class: electronic 0.6854460 0.7525773 0.7174447 0.24742268
## Class: hip-hop    0.6084906 0.6386139 0.6231884 0.36138614
## Class: jazz       0.6956522 0.5685279 0.6256983 0.43147208
## Class: pop        0.5707965 0.6515152 0.6084906 0.34848485
## Class: reggae     0.7201835 0.6977778 0.7088036 0.30222222
## Class: rock       0.7128205 0.7277487 0.7202073 0.27225131

```

3.4 Gradient Boosting

Gradient Boosting is a powerful machine learning technique that builds a predictive model in the form of an ensemble of weak prediction models of decision trees. It sequentially trains new models to correct errors made by existing models, with each new model focusing on the residuals of the previous model. Gradient Boosting is well-suited for regression and classification tasks, especially when dealing with complex relationships and high-dimensional datasets

To optimize the performance of the Gradient Boosting model and prevent overfitting, we employ cross-validation techniques. We use the `trainControl()` function from the “caret” package to split the dataset into training and validation sets. By varying parameters such as the learning rate, tree depth, and the number of boosting iterations, we tune the model for optimal performance.

In this section, we apply Gradient Boosting using the `gbm()` function from the “gbm” package in R. The `gbm()` function iteratively fits a sequence of regression trees to the residuals of the previous model. We evaluate the performance of the Gradient Boosting model using accuracy. Additionally, we analyze the importance of predictor variables using a relative influence measure.

```
library(caret)
library(gbm)
```

```
## Warning: il pacchetto 'gbm' è stato creato con R versione 4.3.2
```

```
## Loaded gbm 2.1.9
```

```
## This version of gbm is no longer under development. Consider transitioning to gbm3, https://github.com
```

```
# Set up parameter grid for tuning (without subsample and colsample_bytree)
param_grid <- expand.grid(
  n.trees = c(50, 100, 500, 1000), # Adjust the values as needed
  interaction.depth = c(1, 2, 4, 6, 8, 10),
  shrinkage = c(0.01, 0.1, 0.3),
  n.minobsinnode = 10
)
param_grid <- expand.grid(
  n.trees = c(1000, 1500, 2000, 5000), # Adjust the values as needed
  interaction.depth = c(1, 5, 10, 15, 20),
  shrinkage = c(0.001, 0.01, 0.1),
  n.minobsinnode = 10
)
```

```
results <- list()
```

```
# Perform grid search over hyperparameter space
for (i in 1:nrow(param_grid)) {
  # Define parameters for this iteration
  params <- list(
    n.trees = param_grid$n.trees[i],
    interaction.depth = param_grid$interaction.depth[i],
    shrinkage = param_grid$shrinkage[i],
    n.minobsinnode = param_grid$n.minobsinnode[i],
    distribution = "multinomial"
  )
}
```

```

# Perform cross-validation with caret and gbm
ctrl <- trainControl(method = "cv", number = 5)
cv_result <- train(
  x = df_filtered[train, -16], # Exclude the response variable column
  y = as.factor(y.train),
  method = "gbm",
  trControl = ctrl,
  tuneGrid = data.frame(n.trees = params$n.trees, interaction.depth = params$interaction.depth, shrinkage = params$shrinkage)
)

# Store the results
results[[paste0("n.trees_", params$n.trees, "_interaction_depth_", params$interaction.depth, "_shrinkage_", params$shrinkage)] <- cv_result$results
}

# Get the best model
#best_index <- which.min(sapply(results, function(x) x$RMSE))
best_model <- results[[which.max(lapply(results, function(x) x$results$Accuracy))]]
print(best_model)

```

The parameter values that we will use for the model are the following:

```

model_parameters <- data.frame(
  n.trees = best_model$bestTune$n.trees,
  interaction.depth = best_model$bestTune$interaction.depth,
  shrinkage = best_model$bestTune$shrinkage,
  n.minobsinnode = best_model$bestTune$n.minobsinnode
)

# Print the dataframe
print(model_parameters)

```

```

##   n.trees interaction.depth shrinkage n.minobsinnode
## 1    2000             10      0.01             10

```

```

final_model <- train(
  x = df_filtered[, -16], # Exclude the response variable column
  y = as.factor(df_filtered$track_genre),
  method = "gbm",
  trControl = trainControl(method = "none"), # No resampling for final training
  tuneGrid = data.frame(
    n.trees = best_model$bestTune$n.trees,
    interaction.depth = best_model$bestTune$interaction.depth,
    shrinkage = best_model$bestTune$shrinkage,
    n.minobsinnode = best_model$bestTune$n.minobsinnode
  )
)

```

```

test_predictions <- predict(final_model, newdata = df_filtered[test, -16])
confusion_matrix <- table(test_predictions, y.test)
accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
print(paste("Accuracy:", accuracy))

```



```
## [1] "Accuracy: 0.9"
```

```
print(confusion_matrix)
```

```
##                y.test
## test_predictions blues classical country electronic hip-hop jazz pop reggae
##      blues      162         0         6         3         0      12     0       0
##      classical    0        206         0         0         0         1     0       0
##      country      1         0        179         0         1         2     2       2
##      electronic   2         0         0        186         1         5     0       1
##      hip.hop      1         0         1         1        160         0    14       8
##      jazz        15         1         2         0         1       174     0       0
##      pop          1         0         2         2         17         2    164       6
##      reggae       1         0         2         1         20         0    14      208
##      rock         6         0         5         1         2         1     4       0
##                y.test
## test_predictions rock
##      blues        1
##      classical     0
##      country       6
##      electronic    0
##      hip.hop       1
##      jazz          0
##      pop           1
##      reggae        1
##      rock         181
```

```
# Print or use the metrics_df data frame
print(metrics_df)
```

```
##                Precision    Recall  F1_Score Specificity
## Class: blues      0.8804348 0.8571429 0.8686327 0.142857143
## Class: classical  0.9951691 0.9951691 0.9951691 0.004830918
## Class: country    0.9274611 0.9086294 0.9179487 0.091370558
## Class: electronic 0.9538462 0.9587629 0.9562982 0.041237113
## Class: hip-hop    0.8602151 0.7920792 0.8247423 0.207920792
## Class: jazz       0.9015544 0.8832487 0.8923077 0.116751269
## Class: pop        0.8410256 0.8282828 0.8346056 0.171717172
## Class: reggae     0.8421053 0.9244444 0.8813559 0.075555556
## Class: rock       0.9050000 0.9476440 0.9258312 0.052356021
```

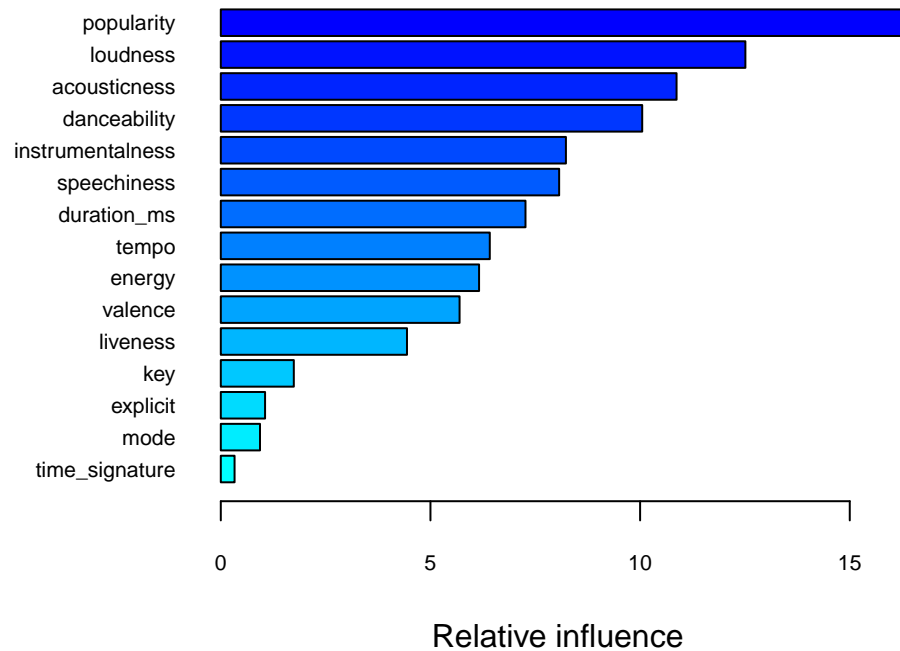
```
mai.old<-par()$mai
mai.old
```

```
## [1] 1.02 0.82 0.82 0.42
```

```
#new vector
mai.new<-mai.old
#new space on the left
mai.new[2] <- 2.5
mai.new
```

```
## [1] 1.02 2.50 0.82 0.42
```

```
#modify graphical parameters  
par(mai=mai.new,cex.axis=0.7)  
#summary(final_model, las=1)  
#las=1 horizontal names on y  
summary(final_model, las=1, cBar=15)
```



```
##           var    rel.inf  
## popularity      popularity 16.2308476  
## loudness        loudness 12.5111344  
## acousticness    acousticness 10.8698713  
## danceability    danceability 10.0507139  
## instrumentalness instrumentalness 8.2308774  
## speechiness     speechiness 8.0696686  
## duration_ms     duration_ms 7.2665409  
## tempo           tempo 6.4153764  
## energy          energy 6.1597559  
## valence         valence 5.6947981  
## liveness        liveness 4.4403517  
## key             key 1.7416439  
## explicit        explicit 1.0555830  
## mode            mode 0.9348006  
## time_signature  time_signature 0.3280362
```

```
#cBar defines how many variables  
#back to original window  
par(mai=mai.old)
```